

## Business Problem Understanding

The COVID-19 pandemic has revealed significant disparities in health outcomes based on patient demographics, preexisting conditions, and vaccination coverage. Public health authorities aim to reduce hospitalization rates, manage reinfection risk, and improve recovery outcomes by identifying key risk factors and evaluating the effectiveness of vaccination programs.

This project seeks to analyze detailed patient-level COVID-19 data to uncover patterns in disease severity, recovery duration, reinfection trends, and long COVID occurrences. The analysis will also help assess how vaccination type, dosage, and patient profiles influence recovery and hospitalization outcomes.

Insights from this analysis will support healthcare decision-makers in optimizing vaccination strategies, allocating medical resources efficiently, and designing targeted interventions for high-risk populations.

## Data Understanding

```
In [13]: import pandas as pd
from scipy.stats import skew
import numpy as np

import seaborn as sns
import matplotlib.pyplot as plt

import warnings
warnings.simplefilter("ignore")
```

```
In [14]: df = pd.read_csv("covid_related_disease_data.csv")
df
```

Out[14]:

	Patient_ID	Age	Gender	Region	Preexisting_Condition	Date_of_Infection
0	1	69	Male	Hovedstaden	Obesity	2022-06-21
1	2	38	Male	Sjælland	Asthma	2024-02-02
2	3	41	Female	Syddanmark	Hypertension	2023-05-28
3	4	81	Female	Hovedstaden	Asthma	2023-08-13
4	5	50	Female	Syddanmark	Cardiovascular	2023-03-10
...	...	...	...	...	...	...
2995	2996	43	Male	Nordjylland	Hypertension	2022-10-19
2996	2997	36	Female	Syddanmark	Obesity	2022-12-16
2997	2998	75	Female	Sjælland	Cardiovascular	2023-09-30
2998	2999	45	Female	Hovedstaden	Asthma	2023-06-06
2999	3000	83	Female	Midtjylland	Obesity	2023-09-07

3000 rows × 26 columns



In [15]:

```
df.shape
```

Out[15]:

(3000, 26)

In [16]:

```
df.size
```

Out[16]:

78000

In [17]:

```
df.head(5)
```

Out[17]:

	Patient_ID	Age	Gender	Region	Preexisting_Condition	Date_of_Infection	COV
0	1	69	Male	Hovedstaden	Obesity	2022-06-21	
1	2	38	Male	Sjælland	Asthma	2024-02-02	
2	3	41	Female	Syddanmark	Hypertension	2023-05-28	
3	4	81	Female	Hovedstaden	Asthma	2023-08-13	
4	5	50	Female	Syddanmark	Cardiovascular	2023-03-10	

5 rows × 26 columns



In [18]:

```
df.tail(3)
```

Out[18]:

	Patient_ID	Age	Gender	Region	Preexisting_Condition	Date_of_Infection
2997	2998	75	Female	Sjælland	Cardiovascular	2023-09-30
2998	2999	45	Female	Hovedstaden	Asthma	2023-06-06
2999	3000	83	Female	Midtjylland	Obesity	2023-09-07

3 rows × 26 columns



In [19]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3000 entries, 0 to 2999
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Patient_ID                            3000 non-null   int64
1   Age                                    3000 non-null   int64
2   Gender                                3000 non-null   object
3   Region                                3000 non-null   object
4   Preexisting_Condition                  2531 non-null   object
5   Date_of_Infection                      3000 non-null   object
6   COVID_Strain                          3000 non-null   object
7   Symptoms                              3000 non-null   object
8   Severity                              3000 non-null   object
9   Hospitalized                          3000 non-null   object
10  Hospital_Admission_Date                876 non-null    object
11  Hospital_Discharge_Date                876 non-null    object
12  ICU_Admission                          3000 non-null   object
13  Ventilator_Support                     3000 non-null   object
14  Recovered                              3000 non-null   object
15  Date_of_Recovery                       1508 non-null   object
16  Reinfection                            3000 non-null   object
17  Date_of_Reinfection                    285 non-null    object
18  Vaccination_Status                     3000 non-null   object
19  Vaccine_Type                           1191 non-null   object
20  Doses_Received                         3000 non-null   int64
21  Date_of_Last_Dose                      1472 non-null   object
22  Long_COVID_Symptoms                    220 non-null    object
23  Occupation                             3000 non-null   object
24  Smoking_Status                         3000 non-null   object
25  BMI                                    3000 non-null   float64
dtypes: float64(1), int64(3), object(22)
memory usage: 609.5+ KB
```

In [20]: `df.columns.tolist()`

```
Out[20]: ['Patient_ID',  
          'Age',  
          'Gender',  
          'Region',  
          'Preexisting_Condition',  
          'Date_of_Infection',  
          'COVID_Strain',  
          'Symptoms',  
          'Severity',  
          'Hospitalized',  
          'Hospital_Admission_Date',  
          'Hospital_Discharge_Date',  
          'ICU_Admission',  
          'Ventilator_Support',  
          'Recovered',  
          'Date_of_Recovery',  
          'Reinfection',  
          'Date_of_Reinfection',  
          'Vaccination_Status',  
          'Vaccine_Type',  
          'Doses_Received',  
          'Date_of_Last_Dose',  
          'Long_COVID_Symptoms',  
          'Occupation',  
          'Smoking_Status',  
          'BMI']
```

```
In [21]: print(df)
```

	Patient_ID	Age	Gender	Region	Preexisting_Condition	\
0	1	69	Male	Hovedstaden	Obesity	
1	2	38	Male	Sjælland	Asthma	
2	3	41	Female	Syddanmark	Hypertension	
3	4	81	Female	Hovedstaden	Asthma	
4	5	50	Female	Syddanmark	Cardiovascular	
...	...	...	...	...	...	
2995	2996	43	Male	Nordjylland	Hypertension	
2996	2997	36	Female	Syddanmark	Obesity	
2997	2998	75	Female	Sjælland	Cardiovascular	
2998	2999	45	Female	Hovedstaden	Asthma	
2999	3000	83	Female	Midtjylland	Obesity	

	Date_of_Infection	COVID_Strain	Symptoms	Severity	Hospitalized	...	\
0	2022-06-21	Delta	Mild	Moderate	Yes	...	
1	2024-02-02	XBB.1.5	Mild	Moderate	No	...	
2	2023-05-28	Beta	Mild	High	Yes	...	
3	2023-08-13	Delta	Severe	High	No	...	
4	2023-03-10	Delta	Mild	High	No	...	
...	...	...	...	...	...	...	
2995	2022-10-19	XBB.1.5	Severe	Critical	No	...	
2996	2022-12-16	Omicron	Moderate	Low	No	...	
2997	2023-09-30	Beta	Severe	Moderate	No	...	
2998	2023-06-06	Delta	Severe	Moderate	No	...	
2999	2023-09-07	XBB.1.5	Moderate	Low	No	...	

	Reinfection	Date_of_Reinfection	Vaccination_Status	Vaccine_Type	\
0	No	NaN	Yes	NaN	
1	No	NaN	No	NaN	
2	No	NaN	Yes	Janssen	
3	Yes	2024-08-24	Yes	AstraZeneca	
4	No	NaN	Yes	NaN	
...	...	...	...	...	
2995	No	NaN	Yes	NaN	
2996	No	NaN	Yes	Pfizer	
2997	No	NaN	Yes	Moderna	
2998	No	NaN	Yes	AstraZeneca	
2999	No	NaN	No	NaN	

	Doses_Received	Date_of_Last_Dose	Long_COVID_Symptoms	Occupation	\
0	1	2022-09-22	NaN	Healthcare	
1	0	NaN	NaN	Healthcare	
2	3	2024-05-14	NaN	Unemployed	
3	1	2024-10-31	NaN	Office Worker	
4	2	2023-07-05	NaN	Student	
...	...	...	...	...	
2995	1	2024-09-20	NaN	Driver	
2996	2	2023-10-05	NaN	Healthcare	
2997	3	2023-05-13	NaN	Teacher	
2998	1	2024-05-13	NaN	Student	
2999	0	NaN	NaN	Teacher	

	Smoking_Status	BMI
0	Never	27.7
1	Never	21.9
2	Never	22.7
3	Never	27.7
4	Never	11.9
...	...	...
2995	Never	22.0

2996	Never	27.8
2997	Former	20.9
2998	Never	19.3
2999	Former	33.0

[3000 rows x 26 columns]

In [22]: `df.dtypes`

```
Out[22]: Patient_ID          int64
Age              int64
Gender           object
Region           object
Preexisting_Condition  object
Date_of_Infection  object
COVID_Strain      object
Symptoms         object
Severity         object
Hospitalized      object
Hospital_Admission_Date  object
Hospital_Discharge_Date  object
ICU_Admission     object
Ventilator_Support  object
Recovered         object
Date_of_Recovery  object
Reinfection       object
Date_of_Reinfection  object
Vaccination_Status  object
Vaccine_Type      object
Doses_Received    int64
Date_of_Last_Dose  object
Long_COVID_Symptoms  object
Occupation        object
Smoking_Status    object
BMI               float64
dtype: object
```

## Data Cleaning / Data Preprocessing

In [25]: `df.columns`

```
Out[25]: Index(['Patient_ID', 'Age', 'Gender', 'Region', 'Preexisting_Condition',
               'Date_of_Infection', 'COVID_Strain', 'Symptoms', 'Severity',
               'Hospitalized', 'Hospital_Admission_Date', 'Hospital_Discharge_Date',
               'ICU_Admission', 'Ventilator_Support', 'Recovered', 'Date_of_Recovery',
               'Reinfection', 'Date_of_Reinfection', 'Vaccination_Status',
               'Vaccine_Type', 'Doses_Received', 'Date_of_Last_Dose',
               'Long_COVID_Symptoms', 'Occupation', 'Smoking_Status', 'BMI'],
              dtype='object')
```

```
In [26]: # to check the duplicated record
df.duplicated().sum()
```

Out[26]: 0

In [27]: `df.shape[0]`

Out[27]: 3000

```
In [28]: # to check the missing values  
df.isnull().sum()
```

```
Out[28]: Patient_ID          0  
Age              0  
Gender           0  
Region          0  
Preexisting_Condition  469  
Date_of_Infection  0  
COVID_Strain     0  
Symptoms         0  
Severity         0  
Hospitalized     0  
Hospital_Admission_Date  2124  
Hospital_Discharge_Date  2124  
ICU_Admission    0  
Ventilator_Support  0  
Recovered        0  
Date_of_Recovery  1492  
Reinfection      0  
Date_of_Reinfection  2715  
Vaccination_Status  0  
Vaccine_Type     1809  
Doses_Received   0  
Date_of_Last_Dose  1528  
Long_COVID_Symptoms  2780  
Occupation       0  
Smoking_Status   0  
BMI              0  
dtype: int64
```

```
In [29]: (df.isnull().sum()/len(df)) * 100
```

```
Out[29]: Patient_ID      0.000000
        Age            0.000000
        Gender         0.000000
        Region         0.000000
        Preexisting_Condition 15.633333
        Date_of_Infection 0.000000
        COVID_Strain     0.000000
        Symptoms        0.000000
        Severity        0.000000
        Hospitalized    0.000000
        Hospital_Admission_Date 70.800000
        Hospital_Discharge_Date 70.800000
        ICU_Admission   0.000000
        Ventilator_Support 0.000000
        Recovered       0.000000
        Date_of_Recovery 49.733333
        Reinfection     0.000000
        Date_of_Reinfection 90.500000
        Vaccination_Status 0.000000
        Vaccine_Type    60.300000
        Doses_Received  0.000000
        Date_of_Last_Dose 50.933333
        Long_COVID_Symptoms 92.666667
        Occupation     0.000000
        Smoking_Status  0.000000
        BMI            0.000000
        dtype: float64
```

```
In [30]: # Check if there are any missing values in the entire DataFrame
print(df.isnull().values.any())
```

True

```
In [31]: # filling missing values insted of nan -> None , means the patient had no prexi
df['Preexisting_Condition'].fillna('None', inplace=True)
```

```
In [34]: # Hospital_Admission_Date and Hospital_Discharge_Date HAVING MORETHAN 30% OF NUL
df.drop(['Hospital_Admission_Date', 'Hospital_Discharge_Date'], axis=1, inplace=
```

```
In [42]: # Replace NaN values in 'Recovery_Days' with the string 'None'
df['Recovery_Days'] = df['Recovery_Days'].fillna(0)
```

```
In [44]: # Step 1: Convert to numeric (non-numeric values become NaN)
df['Recovery_Days'] = pd.to_numeric(df['Recovery_Days'], errors='coerce')
# Step 2: Replace NaN with 0
df['Recovery_Days'] = df['Recovery_Days'].fillna(0)
# Step 3: Convert the column to integer type
df['Recovery_Days'] = df['Recovery_Days'].astype(int)
```

```
In [ ]: # IN Date_of_Recovery having more than 30% of Null Values so i am drop the colu
df.drop(['Date_of_Recovery'], axis=1, inplace=True)
```

```
In [ ]: # drop the Date_of_Reinfection columnin the datasets having more than 30% of nul
df.drop(['Date_of_Reinfection'], axis=1, inplace=True)
```

```
In [ ]: # This Line removes the Date_of_Last_Dose column from the DataFrame df
df.drop(['Date_of_Last_Dose'], axis=1, inplace=True)
```



```

In [ ]: # This line replaces all missing values in the Long_COVID_Symptoms column with t
df['Long_COVID_Symptoms'] = df['Long_COVID_Symptoms'].fillna("None")

In [ ]: df['Vaccine_Type'] = df['Vaccine_Type'].fillna("None")

In [ ]: df.isnull().sum()

In [ ]: # Calculate Q1, Q3, and IQR for the BMI column
Q1 = df['BMI'].quantile(0.25)
Q3 = df['BMI'].quantile(0.75)
IQR = Q3 - Q1

# Define lower and upper bounds for detecting outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

In [ ]: # Remove outliers from the DataFrame based on BMI
df_no_outliers = df[(df['BMI'] >= lower_bound) & (df['BMI'] <= upper_bound)]

In [ ]: # Step 1: Calculate Q1, Q3 and IQR
Q1 = df['Recovery_Days'].quantile(0.25)
Q3 = df['Recovery_Days'].quantile(0.75)
IQR = Q3 - Q1

# Step 2: Define outlier bounds
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Step 3: Calculate median (excluding outliers)
median_value = df[(df['Recovery_Days'] >= lower_bound) & (df['Recovery_Days'] <=
# Step 4: Replace outliers with the median
df['Recovery_Days'] = df['Recovery_Days'].apply(lambda x: median_value if x < lo

# Final output
print(df)

In [ ]: # Display the number of rows before and after removal
original_count = len(df)
cleaned_count = len(df_no_outliers)

original_count, cleaned_count

```

## Feature Engineering

```

In [35]: print(df['Recovered'].unique())
# print(df['Date_of_Recovery'].isna().sum())

['Yes' 'No']

In [36]: # Clean Recovered column
df['Recovered'] = df['Recovered'].fillna('Unknown')
df['Recovered_cleaned'] = df['Recovered'].str.strip().str.lower()

# Create Death column: if not recovered, then assume death
df['Death'] = df['Recovered_cleaned'].apply(lambda x: 'No' if x == 'yes' else 'Y

```

```
# Print how many people are marked as Death = Yes or No
print(df['Death'].value_counts())
```

```
Death
No      1508
Yes     1492
Name: count, dtype: int64
```

```
In [37]: df["Death"].unique()
```

```
Out[37]: array(['No', 'Yes'], dtype=object)
```

```
In [38]: # Show the new Death column with related columns for verification
df[['Patient_ID', 'Recovered', 'Date_of_Recovery', 'Death']].head()
```

```
Out[38]:
```

	Patient_ID	Recovered	Date_of_Recovery	Death
0	1	Yes	2023-04-19	No
1	2	No	NaN	Yes
2	3	No	NaN	Yes
3	4	Yes	2025-02-09	No
4	5	No	NaN	Yes

## Checking the skewness for Numerical columns

```
In [182... # Select numeric columns
numeric_cols = ['Age', 'Doses_Received', 'BMI']

# Calculate skewness for each column
skewness = df[numeric_cols].apply(skew)

skewness
```

```
Out[182... Age      -0.015157
Doses_Received  0.682816
BMI      -0.029625
dtype: float64
```

Age: Skewness  $\approx -0.015$  — distribution is nearly symmetric.(Almost symmetrical)

Doses\_Received: Skewness  $\approx +0.683$  — moderately right-skewed; some individuals received higher doses.(Moderately positively skewed (right-tail))

BMI: Skewness  $\approx -0.030$  — distribution is almost symmetric with slight left tilt.(Almost symmetrical)


## log Transformation

```
In [183... # Apply Log transformation to Doses_Received
df['Doses_Received_Log'] = np.log1p(df['Doses_Received']) # Handles 0 values sa
```

```
In [184... # Plot before and after transformation
plt.figure(figsize=(12, 5))

# Original
plt.subplot(1, 2, 1)
sns.histplot(df['Doses_Received'], kde=True, color='green')
plt.title("Original Doses_Received")

# Transformed
plt.subplot(1, 2, 2)
sns.histplot(df['Doses_Received_Log'], kde=True, color='red')
plt.title("Log Transformed Doses_Received")
plt.tight_layout()
plt.show()
```



The original Doses\_Received column was moderately right-skewed with a long tail of high values. After applying log transformation, the distribution became more symmetric and compact, making it suitable for modeling.

```
In [185... # Age Grouping
df['Age_Group'] = pd.cut(df['Age'], bins=[0, 18, 35, 50, 65, 100], labels=['Child', 'Young Adult', 'Adult', 'Older Adult', 'Elderly'])
```

```
In [186... # BMI Category
df['BMI_Category'] = pd.cut(df['BMI'], bins=[0, 18.5, 24.9, 29.9, 100], labels=['Underweight', 'Normal Weight', 'Overweight', 'Obese'])
```

```
In [187... # Total Risk Score (mock feature combining Age and BMI)
df['Risk_Score'] = df['Age'] * df['BMI']
```

```
In [41]: # Recovery Time
if 'Date_of_Recovery' in df.columns and 'Date_of_Infection' in df.columns:
    df['Date_of_Recovery'] = pd.to_datetime(df['Date_of_Recovery'], errors='coerce')
    df['Date_of_Infection'] = pd.to_datetime(df['Date_of_Infection'], errors='coerce')
    df['Recovery_Days'] = (df['Date_of_Recovery'] - df['Date_of_Infection']).dt.days
```

```
In [189... # Preview new columns
print(df[['Age_Group', 'BMI_Category', 'Risk_Score', 'Recovery_Days']].head())
```

	Age_Group	BMI_Category	Risk_Score	Recovery_Days
0	Elderly	Overweight	1911.3	302.0
1	Adult	Normal	832.2	NaN
2	Adult	Normal	930.7	NaN
3	Elderly	Overweight	2243.7	546.0
4	Adult	Underweight	595.0	NaN

```
In [190... # Show only numeric columns
numeric_cols = df.select_dtypes(include='number')

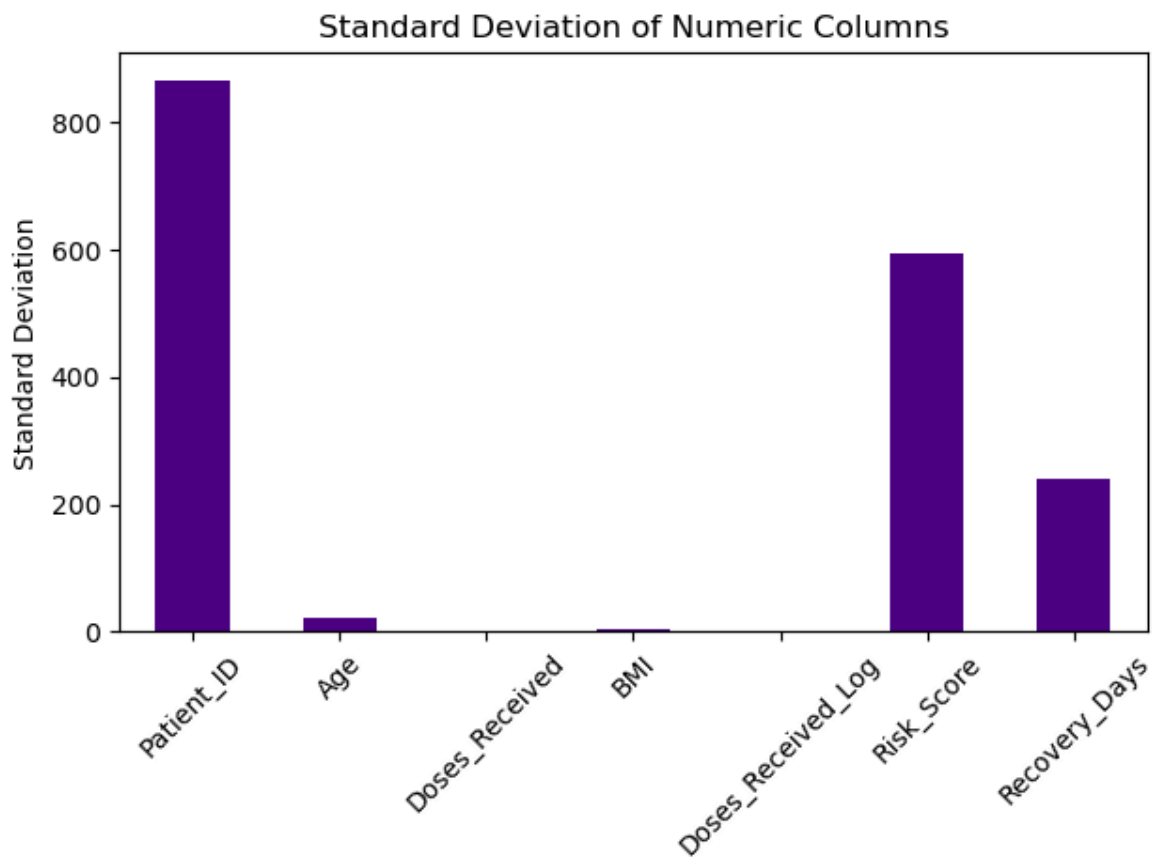
# Calculate standard deviation
std_devs = numeric_cols.std()

# Display
print(f"📊 Standard Deviation of Numeric Columns:\n")
print(std_devs)
```

📊 Standard Deviation of Numeric Columns:

Patient_ID	866.169729
Age	20.872919
Doses_Received	1.154025
BMI	4.898435
Doses_Received_Log	0.566189
Risk_Score	593.294407
Recovery_Days	239.422352
dtype:	float64

```
In [191... std_devs.plot(kind='bar', color='indigo')
plt.title('Standard Deviation of Numeric Columns')
plt.ylabel('Standard Deviation')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Visualise the standard deviation using plots

## EDA (Exploratory Data Analysis)

Seprate the each n every column as per given data

```
In [192... continuous_cols = ["Age", "Recovery_Days", "BMI", "Risk_Score"]

discrete_cols = ["Doses_Received"]

categorical_cols = ["Gender", "Age_Group", "BMI_Category", "Long_COVID_Symptoms", "
```

Count the unique values and values counts for each columns

```
In [193... for i in continuous_cols:
    print(i,':',df[i].unique())
    print('=====')
    print(i,':',df[i].nunique())
    print('=====')
    print(i,':',df[i].value_counts())
```

Age : [69 38 41 81 50 66 76 77 79 72 20 56 35 70 64 53 23 71 80 61 89 32 82 18  
 25 22 45 40 54 52 44 68 59 87 49 75 73 19 85 34 55 28 33 36 84 88 74 26  
 37 42 65 39 86 63 43 24 27 21 67 29 60 30 51 46 47 62 83 31 57 48 58 78]

=====  
 Age : 72  
 =====

=====  
 Age : Age

85 55  
 30 52  
 75 52  
 71 51  
 79 51  
 ..  
 65 34  
 58 34  
 78 31  
 24 30  
 67 28

Name: count, Length: 72, dtype: int64

Recovery\_Days : [3.020e+02 nan 5.460e+02 6.200e+02 2.390e+02 6.190e+02 6.10  
 0e+01

9.100e+01 3.350e+02 4.090e+02 1.040e+02 4.200e+01 2.020e+02 2.060e+02  
 4.570e+02 5.810e+02 5.090e+02 4.560e+02 3.900e+02 9.750e+02 1.640e+02  
 4.000e+01 7.540e+02 3.650e+02 8.510e+02 3.150e+02 2.080e+02 1.440e+02  
 3.600e+02 5.000e+00 2.280e+02 2.610e+02 3.810e+02 7.000e+00 1.800e+02  
 3.630e+02 4.260e+02 4.930e+02 4.290e+02 6.550e+02 2.090e+02 2.230e+02  
 2.640e+02 6.830e+02 9.900e+01 1.010e+02 3.740e+02 1.540e+02 2.050e+02  
 1.610e+02 1.930e+02 4.770e+02 3.270e+02 2.750e+02 5.250e+02 1.330e+02  
 4.370e+02 4.200e+02 3.440e+02 8.810e+02 1.980e+02 3.300e+02 5.340e+02  
 5.650e+02 1.000e+02 8.630e+02 6.710e+02 3.250e+02 4.720e+02 8.800e+01  
 1.150e+02 7.280e+02 8.430e+02 2.140e+02 2.000e+00 1.730e+02 7.800e+01  
 1.900e+02 2.620e+02 5.060e+02 6.610e+02 1.000e+00 5.530e+02 2.600e+01  
 2.220e+02 5.830e+02 7.980e+02 2.580e+02 4.680e+02 3.060e+02 3.140e+02  
 2.250e+02 1.520e+02 1.550e+02 6.170e+02 2.700e+02 5.800e+01 5.550e+02  
 4.960e+02 3.320e+02 5.110e+02 1.680e+02 4.220e+02 2.550e+02 5.200e+01  
 8.460e+02 2.540e+02 4.100e+01 4.920e+02 3.800e+02 1.000e+01 8.100e+01  
 1.380e+02 2.480e+02 1.200e+01 2.210e+02 2.010e+02 3.100e+02 7.830e+02  
 3.560e+02 6.500e+01 7.100e+01 2.780e+02 4.800e+01 2.900e+02 2.430e+02  
 8.080e+02 2.100e+01 5.790e+02 7.640e+02 2.310e+02 8.760e+02 6.960e+02  
 7.600e+02 6.730e+02 2.110e+02 1.590e+02 3.820e+02 2.720e+02 4.410e+02  
 2.300e+02 3.770e+02 1.140e+02 7.560e+02 6.460e+02 6.880e+02 5.670e+02  
 3.700e+01 8.500e+01 2.200e+02 3.310e+02 7.000e+01 4.270e+02 3.370e+02  
 1.770e+02 3.360e+02 1.620e+02 7.400e+02 3.500e+02 3.080e+02 6.530e+02  
 2.040e+02 3.130e+02 8.200e+01 4.490e+02 4.970e+02 7.420e+02 3.070e+02  
 1.400e+01 3.040e+02 6.260e+02 6.240e+02 5.600e+01 1.230e+02 9.600e+01  
 3.900e+01 4.690e+02 3.680e+02 4.760e+02 6.750e+02 2.590e+02 7.550e+02  
 9.500e+01 1.190e+02 7.570e+02 4.520e+02 2.880e+02 5.660e+02 3.490e+02  
 3.750e+02 2.930e+02 3.450e+02 1.960e+02 1.530e+02 4.080e+02 1.750e+02  
 7.790e+02 2.950e+02 7.300e+01 5.600e+02 8.600e+02 1.670e+02 5.470e+02  
 4.670e+02 5.450e+02 6.630e+02 9.060e+02 9.110e+02 6.000e+00 5.260e+02  
 3.380e+02 7.400e+01 1.490e+02 7.470e+02 5.890e+02 1.710e+02 7.900e+01  
 5.040e+02 4.750e+02 3.600e+01 5.570e+02 7.500e+02 3.390e+02 1.320e+02  
 1.470e+02 3.880e+02 7.670e+02 9.800e+01 1.210e+02 5.400e+01 3.410e+02  
 5.350e+02 3.920e+02 4.000e+00 4.430e+02 4.340e+02 1.510e+02 8.300e+01  
 8.100e+02 9.800e+02 9.130e+02 2.670e+02 3.620e+02 1.310e+02 6.450e+02  
 1.360e+02 4.000e+02 6.640e+02 3.670e+02 3.730e+02 3.720e+02 4.890e+02  
 5.500e+01 2.850e+02 4.860e+02 1.870e+02 4.350e+02 8.230e+02 1.077e+03  
 3.470e+02 1.120e+02 1.890e+02 1.180e+02 6.780e+02 8.870e+02 4.710e+02

4.870e+02	3.180e+02	5.900e+02	3.700e+02	5.780e+02	1.760e+02	5.240e+02
5.230e+02	4.730e+02	6.220e+02	6.940e+02	3.890e+02	4.280e+02	4.780e+02
4.480e+02	8.380e+02	2.650e+02	6.030e+02	4.440e+02	1.500e+01	3.430e+02
2.400e+01	4.190e+02	1.250e+02	7.380e+02	2.830e+02	3.570e+02	2.380e+02
1.560e+02	6.500e+02	8.160e+02	7.000e+02	1.660e+02	6.380e+02	8.670e+02
3.160e+02	1.350e+02	0.000e+00	3.460e+02	2.180e+02	3.090e+02	1.130e+02
6.180e+02	3.000e+02	9.510e+02	7.200e+01	1.280e+02	8.300e+02	4.580e+02
5.940e+02	9.000e+01	1.060e+02	3.200e+02	2.820e+02	5.160e+02	4.830e+02
6.200e+01	5.760e+02	1.920e+02	2.240e+02	4.130e+02	1.420e+02	8.390e+02
6.270e+02	3.790e+02	9.180e+02	3.520e+02	5.380e+02	4.500e+02	5.170e+02
3.000e+00	5.440e+02	2.360e+02	5.200e+02	4.550e+02	1.370e+02	5.070e+02
1.240e+02	5.080e+02	1.600e+02	6.050e+02	3.530e+02	8.260e+02	1.990e+02
6.400e+01	5.100e+01	9.330e+02	1.460e+02	8.150e+02	5.390e+02	3.580e+02
5.990e+02	1.800e+01	5.730e+02	1.570e+02	4.640e+02	3.170e+02	3.660e+02
4.950e+02	3.330e+02	2.630e+02	3.780e+02	7.300e+02	6.860e+02	8.540e+02
1.400e+02	2.690e+02	1.580e+02	4.250e+02	8.660e+02	4.790e+02	5.140e+02
1.030e+02	4.990e+02	4.360e+02	4.320e+02	3.220e+02	5.860e+02	3.300e+01
1.050e+02	2.530e+02	2.410e+02	5.630e+02	6.690e+02	6.110e+02	5.270e+02
5.820e+02	1.700e+02	1.700e+01	3.610e+02	1.430e+02	6.900e+01	4.380e+02
1.100e+01	9.400e+01	1.830e+02	3.840e+02	6.480e+02	4.600e+02	5.310e+02
2.660e+02	9.480e+02	3.010e+02	2.560e+02	2.330e+02	3.690e+02	3.830e+02
2.130e+02	8.070e+02	9.740e+02	5.120e+02	4.150e+02	6.950e+02	4.310e+02
5.610e+02	5.800e+02	5.000e+02	4.450e+02	2.290e+02	3.190e+02	6.600e+01
5.700e+01	1.600e+01	9.640e+02	1.500e+02	8.190e+02	2.490e+02	2.900e+01
2.910e+02	2.770e+02	5.330e+02	4.110e+02	9.360e+02	2.070e+02	4.040e+02
2.730e+02	6.130e+02	7.070e+02	2.470e+02	4.850e+02	7.140e+02	2.270e+02
2.120e+02	1.690e+02	6.970e+02	3.230e+02	4.900e+02	2.760e+02	9.400e+02
3.260e+02	7.890e+02	8.000e+00	5.850e+02	2.920e+02	2.570e+02	3.850e+02
4.030e+02	4.540e+02	5.920e+02	4.210e+02	1.044e+03	9.030e+02	4.300e+02
6.000e+01	1.042e+03	7.960e+02	1.900e+01	2.970e+02	6.700e+01	5.490e+02
5.410e+02	4.800e+02	3.200e+01	4.600e+01	5.190e+02	3.400e+02	8.480e+02
3.050e+02	6.740e+02	8.270e+02	9.790e+02	1.008e+03	7.770e+02	7.800e+02
2.890e+02	4.660e+02	3.030e+02	5.370e+02	1.340e+02	2.500e+01	9.020e+02
3.590e+02	2.940e+02	4.180e+02	4.050e+02	2.200e+01	5.010e+02	2.710e+02
6.020e+02	1.850e+02	5.950e+02	6.310e+02	8.850e+02	1.790e+02	7.840e+02
1.300e+02	4.300e+01	2.440e+02	4.160e+02	8.470e+02	2.100e+02	4.910e+02
2.860e+02	4.980e+02	4.400e+01	8.550e+02	7.210e+02	5.750e+02	4.630e+02
6.010e+02	2.320e+02	2.520e+02	7.020e+02	3.100e+01	2.160e+02	3.980e+02
9.140e+02	3.510e+02	1.840e+02	2.800e+01	6.470e+02	6.850e+02	3.340e+02
7.430e+02	4.820e+02	7.060e+02	8.730e+02	9.580e+02	2.450e+02	7.700e+02
5.360e+02	1.026e+03	2.370e+02	1.260e+02	5.180e+02	3.940e+02	3.500e+01
6.300e+01	7.760e+02	8.700e+01	5.210e+02	8.900e+01	8.350e+02	7.090e+02
5.910e+02	5.430e+02	6.840e+02	9.690e+02	2.680e+02	7.500e+01	1.160e+02
3.110e+02	5.640e+02	3.420e+02	9.700e+02	5.300e+02	3.960e+02	3.930e+02
7.190e+02	1.051e+03	3.480e+02	4.620e+02	1.950e+02	8.110e+02	5.020e+02
4.700e+02	6.920e+02	8.210e+02	2.500e+02	8.000e+01	2.300e+01	4.530e+02
9.300e+01	4.170e+02	2.700e+01	5.030e+02	3.760e+02	6.900e+02	1.910e+02
5.870e+02	4.330e+02	9.290e+02	6.100e+02	7.200e+02	8.600e+01	7.650e+02
1.810e+02	2.990e+02	7.010e+02	6.250e+02	7.970e+02	9.650e+02	1.860e+02
2.980e+02	7.690e+02	1.080e+02	1.300e+01	5.720e+02	2.340e+02	8.710e+02
6.620e+02	3.910e+02	7.590e+02	5.130e+02	6.670e+02	2.800e+02	3.710e+02
4.390e+02	2.960e+02	5.690e+02	1.020e+02	3.870e+02	7.270e+02	6.390e+02
6.890e+02	2.000e+02	4.470e+02	2.030e+02	7.170e+02	9.710e+02	6.060e+02
6.590e+02	6.650e+02	4.460e+02	4.070e+02	7.160e+02	7.290e+02	1.270e+02
1.170e+02	6.080e+02	9.200e+01	9.920e+02	3.550e+02	9.000e+00	6.980e+02
4.900e+01	1.480e+02	1.063e+03	2.790e+02	9.700e+01	6.820e+02	8.490e+02
1.410e+02	7.700e+01	3.800e+01	1.220e+02	2.510e+02	6.400e+02	8.220e+02
4.650e+02	3.000e+01	3.290e+02	9.590e+02	1.200e+02	1.058e+03	6.700e+02
4.100e+02	4.500e+01	3.400e+01	8.680e+02	6.800e+01	8.030e+02	5.300e+01
7.480e+02	1.940e+02	9.100e+02	2.840e+02	7.410e+02	8.330e+02	8.900e+02

```
2.260e+02 1.450e+02 5.590e+02 8.860e+02 7.080e+02 4.020e+02 5.480e+02
6.370e+02 1.005e+03 1.970e+02 2.460e+02 9.600e+02 1.034e+03 2.000e+01
5.840e+02 6.930e+02 4.140e+02 1.100e+02 8.050e+02 6.520e+02 6.870e+02]
```

```
Recovery_Days : 706
```

```
Recovery_Days : Recovery_Days
```

```
437.0    7
323.0    7
457.0    6
177.0    6
486.0    6
..
431.0    1
695.0    1
415.0    1
974.0    1
687.0    1
```

```
Name: count, Length: 706, dtype: int64
```

```
BMI : [27.7 21.9 22.7 11.9 29.8 22.3 24.4 26.1 21.2 27.1 29.2 22.  29.7 24.9
19.7 18.7 30.5 18.4 26.  28.7 25.9 17.6 21.4 19.3 25.7 24.8 23.6 30.2
28.  20.5 17.1 28.2 26.8 21.8 32.5 26.7 24.2 20.2 29.5 25.1 33.1 29.1
23.4 22.9 18.  11.7 20.9 21.7 24.3 27.3 24.7 25.  23.9 21.6 32.  22.5
29.6 33.5 29.4 20.4 32.7 23.2 27.8 25.3 27.5 31.5 20.7 28.1 31.1 18.9
23.8 38.7 25.5 38.9 15.3 20.3 18.8 19.6 31.9 28.3 17.8 19.1 23.  25.4
15.6 24.5 22.6 21.1 33.3 17.3 23.3 35.6 34.4 14.5 28.8 25.6 35.7 15.2
30.3 32.8 20.1 32.3 26.9 22.8 12.7 24.6 26.6 21.  26.2 34.1 27.  30.8
34.3 32.2 15.4 17.9 12.2 30.7 25.2 21.5 35.9 14.4 22.4 32.4 30.9 34.6
32.6 21.3 18.2 31.7 28.4 31.4 29.  30.  18.6 18.5 14.3 27.6 19.2 34.9
29.3 27.4 20.6 17.2 14.6 20.8 19.9 27.9 24.  32.9 26.5 13.6 26.3 32.1
16.8 26.4 16.6 23.7 27.2 33.  28.9 13.4 19.4 11.5 25.8 22.2 30.6 23.5
28.6 17.5 18.3 16.2 29.9 33.9 30.1 14.9 19.8 19.  22.1 23.1 35.1 31.6
37.6 16.7 30.4 31.3 31.8 31.  24.1 35.  10.8 35.4 40.7 19.5 34.2 33.2
31.2 28.5 34.7 13.3 13.5 12.4 36.4 12.  35.8 33.6 16.1 34.  16.4 36.1
20.  18.1 33.4 11.4 10.4 11.2 37.7 16.9 35.2 35.5 15.7 16.5 15.9 16.
15.5 17.  15.  16.3 36.2 14.2 11.8 42.5 17.7 15.8 37.1 36.5 39.6 34.8
15.1 35.3 33.7 33.8 13.8 34.5 37.8 12.3 17.4 44.6 14.  14.1 10.2 13.
14.7 38.2 12.6 38.4 37.  12.9 13.1 38.8 12.5 11.1 36.9 38.  14.8]
```

```
BMI : 265
```

```
BMI : BMI
```

```
25.6    38
26.7    33
25.7    33
22.9    33
24.9    32
..
11.2     1
36.2     1
11.8     1
42.5     1
14.8     1
```

```
Name: count, Length: 265, dtype: int64
```

```
Risk_Score : [1911.3  832.2  930.7  ... 1567.5  868.5 2739. ]
```



```

=====
Risk_Score : 2436
=====
=====
Risk_Score : Risk_Score
1632.0      6
1320.0      5
858.0       5
756.0       4
1134.0      4
..
1131.6      1
1407.0      1
2705.6      1
888.0       1
2739.0      1
Name: count, Length: 2436, dtype: int64

```

```

In [194... for i in discrete_cols:
            print(i,':',df[i].unique())
            print('=====')
            print(i,':',df[i].nunique())
            print('=====')
            print(i,':',df[i].value_counts())

```

```

Doses_Received : [1 0 3 2]
=====
=====
Doses_Received : 4
=====
=====
Doses_Received : Doses_Received
0      1528
3       497
1       496
2       479
Name: count, dtype: int64

```

```

In [195... for i in categorical_cols:
            print(i,':',df[i].unique())
            print('=====')
            print(i,':',df[i].nunique())
            print('=====')
            print(i,':',df[i].value_counts())

```

```

Gender : ['Male' 'Female']
=====
=====
Gender : 2
=====
=====
Gender : Gender
Female      1527
Male        1473
Name: count, dtype: int64
Age_Group : ['Elderly', 'Adult', 'Youth', 'Senior', 'Child']
Categories (5, object): ['Child' < 'Youth' < 'Adult' < 'Senior' < 'Elderly']
=====
=====
Age_Group : 5
=====
=====
Age_Group : Age_Group
Elderly      1048
Youth         685
Adult         636
Senior        588
Child          43
Name: count, dtype: int64
BMI_Category : ['Overweight', 'Normal', 'Underweight', 'Obese']
Categories (4, object): ['Underweight' < 'Normal' < 'Overweight' < 'Obese']
=====
=====
BMI_Category : 4
=====
=====
BMI_Category : BMI_Category
Normal          1188
Overweight      1058
Obese            488
Underweight      266
Name: count, dtype: int64
Long_COVID_Symptoms : ['None' 'Fatigue' 'Chest Pain' 'Shortness of Breath' 'Brain Fog']
=====
=====
Long_COVID_Symptoms : 5
=====
=====
Long_COVID_Symptoms : Long_COVID_Symptoms
None                2780
Fatigue              62
Brain Fog            59
Chest Pain           52
Shortness of Breath  47
Name: count, dtype: int64
Death : ['No' 'Yes']
=====
=====
Death : 2
=====
=====
Death : Death
No          1508
Yes         1492

```

Name: count, dtype: int64  
 Vaccination\_Status : ['Yes' 'No']

Vaccination\_Status : 2

Vaccination\_Status : Vaccination\_Status

No 1528

Yes 1472

Name: count, dtype: int64

In [196...

```
# for Numerical Variables
df[continuous_cols].describe()
```

Out[196...

	Age	Recovery_Days	BMI	Risk_Score
<b>count</b>	3000.000000	1508.000000	3000.000000	3000.000000
<b>mean</b>	53.944000	362.282493	25.096500	1351.793367
<b>std</b>	20.872919	239.422352	4.898435	593.294407
<b>min</b>	18.000000	0.000000	10.200000	205.200000
<b>25%</b>	36.000000	172.500000	21.800000	859.650000
<b>50%</b>	54.000000	338.000000	25.100000	1298.000000
<b>75%</b>	72.000000	512.000000	28.500000	1773.375000
<b>max</b>	89.000000	1077.000000	44.600000	3570.000000

In [197...

```
# for discrete Variables
df[discrete_cols].describe()
```

Out[197...

	Doses_Received
<b>count</b>	3000.000000
<b>mean</b>	0.981667
<b>std</b>	1.154025
<b>min</b>	0.000000
<b>25%</b>	0.000000
<b>50%</b>	0.000000
<b>75%</b>	2.000000
<b>max</b>	3.000000

In [198...

```
# for categorical Variables
df[categorical_cols].describe()
```

Out[198...

	Gender	Age_Group	BMI_Category	Long_COVID_Symptoms	Death	Vaccination
<b>count</b>	3000	3000	3000	3000	3000	
<b>unique</b>	2	5	4	5	2	
<b>top</b>	Female	Elderly	Normal	None	No	
<b>freq</b>	1527	1048	1188	2780	1508	




## Visualize The Plots

## Univariate Plots

### Checking Outliers For Continuous Column

In [199...

```
sns.boxplot(df["Age"])
plt.show()
```




The boxplot shows that most patients are aged between ~37 and ~73, with a median age of ~55, ranging from ~18 to ~90, and no visible outliers.

In [200...

```
sns.histplot(df['Age'], kde=True, color='blue')
plt.title('Distribution of Age')
```


Out[200...

```
Text(0.5, 1.0, 'Distribution of Age')
```



The age distribution is fairly uniform with slight peaks around ages 30–40 and 80–90, indicating a balanced spread across age groups.

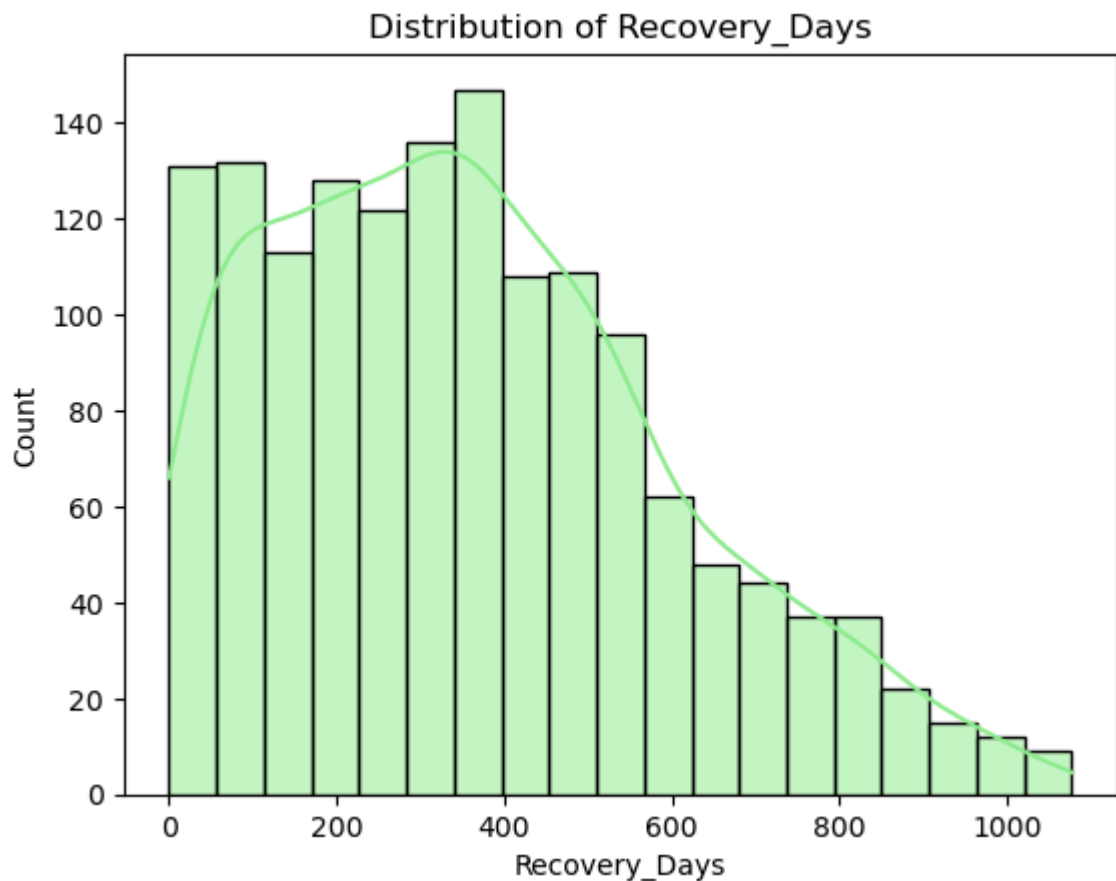
```
In [201... sns.boxplot(df["Recovery_Days"])  
plt.show()
```



The boxplot shows that most patients recovered within 0–350 days, but there are many extreme outliers above 850 days, indicating unusually long recovery times for some cases.


```
In [202... sns.histplot(df['Recovery_Days'], kde=True, color='lightgreen')  
plt.title('Distribution of Recovery_Days')
```

```
Out[202... Text(0.5, 1.0, 'Distribution of Recovery_Days')
```



The distribution of Recovery\_Days is highly right-skewed, indicating that most patients recovered quickly, but a few cases had extremely long recovery durations (potential outliers).


```
In [203... sns.boxplot(df["BMI"])  
plt.show()
```



The BMI boxplot shows that most values lie between 22 and 30, but there are several outliers below 15 and above 38, indicating a few underweight and obese individuals.

```
In [204... sns.histplot(df['BMI'], kde=True, color='violet')
plt.title('Distribution of BMI')
```

```
Out[204... Text(0.5, 1.0, 'Distribution of BMI')
```




The distribution of BMI is slightly right-skewed, indicating that while most individuals have a healthy to slightly overweight BMI, a few outliers have significantly higher values. Most patients have BMI values between 20 and 30, indicating many are in the normal to overweight range.

### CountPlot

```
In [205... # Gender - CountPlot
sns.countplot(data=df, x='Gender', palette='Set2')
plt.title('Count of Gender')
```

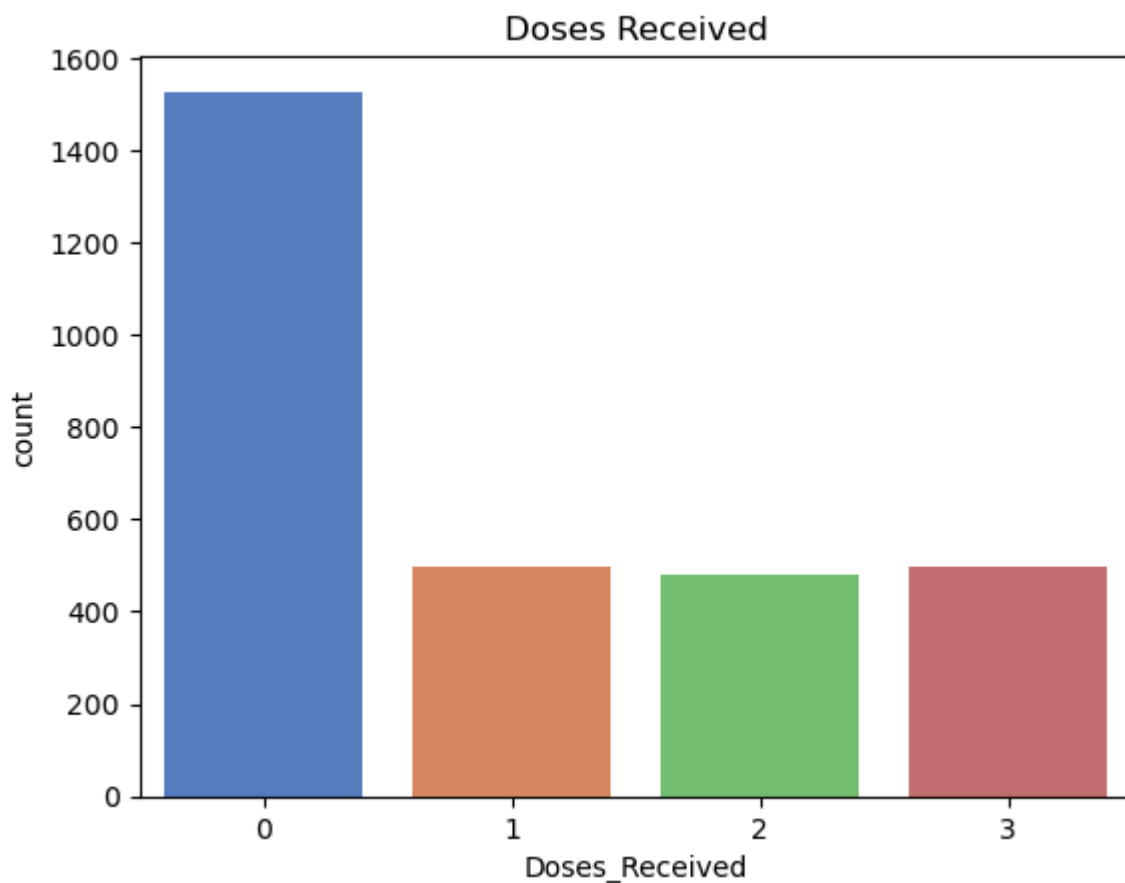
```
Out[205... Text(0.5, 1.0, 'Count of Gender')
```






```
In [206... # Doses Received - Countplot  
sns.countplot(data=df, x='Doses_Received', palette='muted')  
plt.title('Doses Received')
```

```
Out[206... Text(0.5, 1.0, 'Doses Received')
```




```
In [207... # Vaccination Status - Countplot
sns.countplot(data=df, x='Vaccination_Status', palette='coolwarm')
plt.title('Vaccination Status Count')
```

```
Out[207... Text(0.5, 1.0, 'Vaccination Status Count')
```



```
In [208... # Long COVID Symptoms - Countplot
sns.countplot(data=df, x='Long_COVID_Symptoms', palette='Set2')
plt.title('Long COVID Symptoms Presence')
```

```
Out[208... Text(0.5, 1.0, 'Long COVID Symptoms Presence')
```

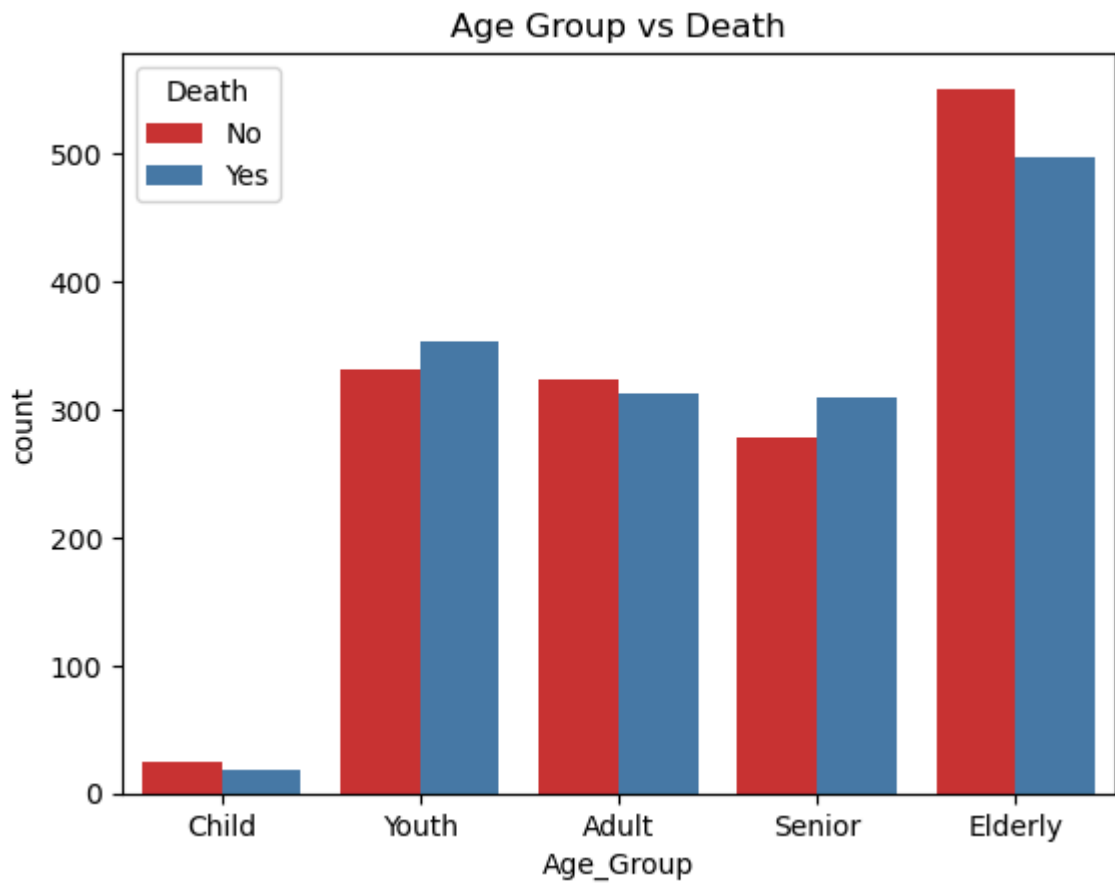


## Bivariate Plots

### Age Group vs Death

```
In [209... # Age Group vs Death
sns.countplot(data=df, x='Age_Group', hue='Death', palette='Set1')
plt.title('Age Group vs Death')
```

```
Out[209... Text(0.5, 1.0, 'Age Group vs Death')
```



## Vaccination Status vs Death

```
In [4]: # Vaccination Status vs Death
sns.countplot(data=df, x='Vaccination_Status', hue='Death', palette='Set2')
plt.title('Vaccination Status vs Death')
```

```

-----
ValueError                                Traceback (most recent call last)
Cell In[4], line 2
      1 # Vaccination Status vs Death
----> 2 sns.countplot(data=df, x='Vaccination_Status', hue='Death', palette='Set
      2')
      3 plt.title('Vaccination Status vs Death')

File ~\anaconda3\Lib\site-packages\seaborn\categorical.py:2631, in countplot(data, x, y, hue, order, hue_order, orient, color, palette, saturation, fill, hue_norm, stat, width, dodge, gap, log_scale, native_scale, formatter, legend, ax, **kwargs)
    2628 elif x is not None and y is not None:
    2629     raise TypeError("Cannot pass values for both `x` and `y`.")
-> 2631 p = _CategoricalAggPlotter(
    2632     data=data,
    2633     variables=dict(x=x, y=y, hue=hue),
    2634     order=order,
    2635     orient=orient,
    2636     color=color,
    2637     legend=legend,
    2638 )
    2640 if ax is None:
    2641     ax = plt.gca()

File ~\anaconda3\Lib\site-packages\seaborn\categorical.py:67, in _CategoricalPlotter.__init__(self, data, variables, order, orient, require_numeric, color, legend)
     56 def __init__(
     57     self,
     58     data=None,
     (... )
     64     legend="auto",
     65 ):
--> 67     super().__init__(data=data, variables=variables)
     69     # This method takes care of some bookkeeping that is necessary because the
     70     # original categorical plots (prior to the 2021 refactor) had some rules that
     71     # don't fit exactly into VectorPlotter logic. It may be wise to have a second
     (... )
     76     # default VectorPlotter rules. If we do decide to make orient part of the
     77     # _base variable assignment, we'll want to figure out how to express that.
     78     if self.input_format == "wide" and orient in ["h", "y"]:

File ~\anaconda3\Lib\site-packages\seaborn\_base.py:634, in VectorPlotter.__init__(self, data, variables)
    629 # var_ordered is relevant only for categorical axis variables, and may
    630 # be better handled by an internal axis information object that tracks
    631 # such information and is set up by the scale_* methods. The analogous
    632 # information for numeric axes would be information about log scales.
    633 self._var_ordered = {"x": False, "y": False} # alt., used DefaultDict
--> 634 self.assign_variables(data, variables)
    636 # TODO Lots of tests assume that these are called to initialize the
    637 # mappings to default values on class initialization. I'd prefer to
    638 # move away from that and only have a mapping when explicitly called.
    639 for var in ["hue", "size", "style"]:

```

```

File ~\anaconda3\Lib\site-packages\seaborn\_base.py:679, in VectorPlotter.assign_variables(self, data, variables)
    674 else:
    675     # When dealing with long-form input, use the newer PlotData
    676     # object (internal but introduced for the objects interface)
    677     # to centralize / standardize data consumption logic.
    678     self.input_format = "long"
--> 679     plot_data = PlotData(data, variables)
    680     frame = plot_data.frame
    681     names = plot_data.names

File ~\anaconda3\Lib\site-packages\seaborn\_core\data.py:58, in PlotData.__init__(self, data, variables)
    51 def __init__(
    52     self,
    53     data: DataSource,
    54     variables: dict[str, VariableSpec],
    55 ):
    57     data = handle_data_source(data)
---> 58     frame, names, ids = self._assign_variables(data, variables)
    60     self.frame = frame
    61     self.names = names

File ~\anaconda3\Lib\site-packages\seaborn\_core\data.py:232, in PlotData._assign_variables(self, data, variables)
    230     else:
    231         err += "An entry with this name does not appear in `data`."
--> 232     raise ValueError(err)
    234 else:
    235
    236     # Otherwise, assume the value somehow represents data
    237
    238     # Ignore empty data structures
    239     if isinstance(val, Sized) and len(val) == 0:

ValueError: Could not interpret value `Death` for `hue`. An entry with this name
does not appear in `data`.

```

## Severity vs Hospitalized

```

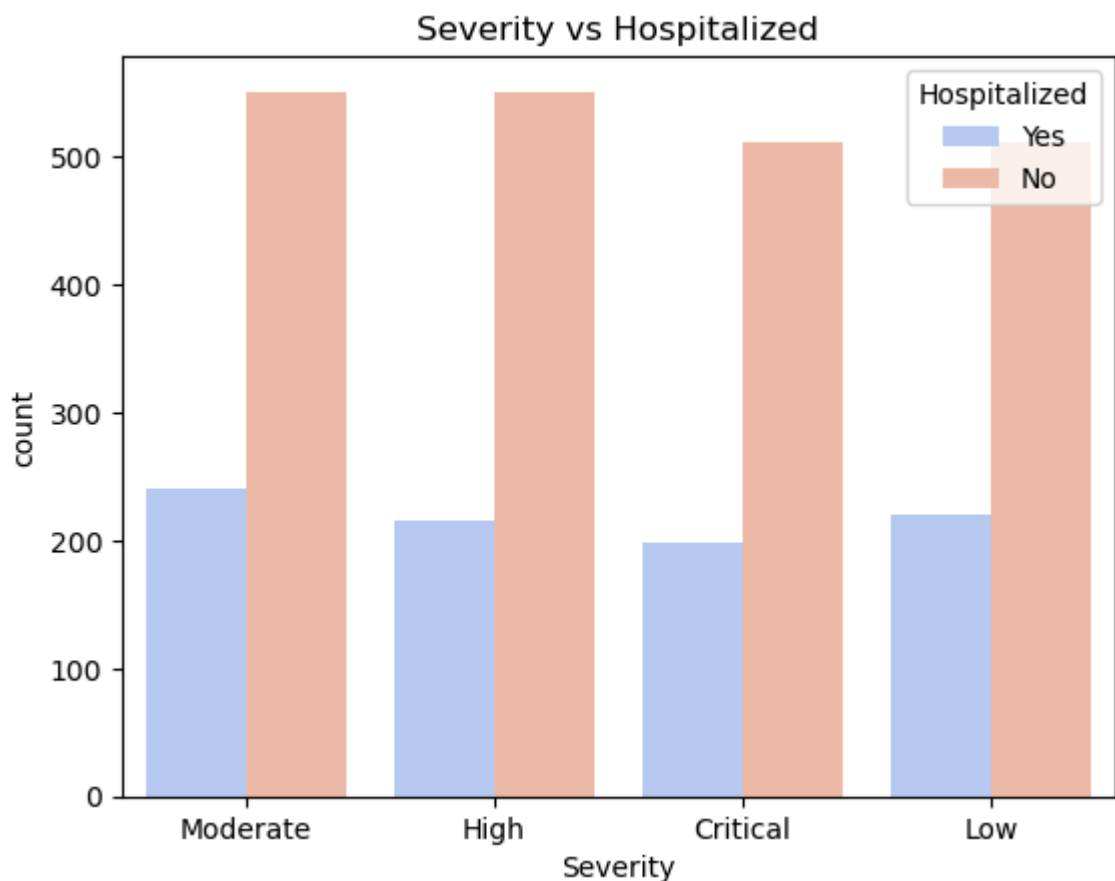
In [211... # Severity vs Hospitalized
sns.countplot(data=df, x='Severity', hue='Hospitalized', palette='coolwarm')
plt.title('Severity vs Hospitalized')

```

```

Out[211... Text(0.5, 1.0, 'Severity vs Hospitalized')


```



## ICU Admission vs Ventilator Support

```
In [212... # ICU Admission vs Ventilator Support
sns.countplot(data=df, x='ICU_Admission', hue='Ventilator_Support', palette='pas
plt.title('ICU Admission vs Ventilator Support')
```

```
Out[212... Text(0.5, 1.0, 'ICU Admission vs Ventilator Support')
```



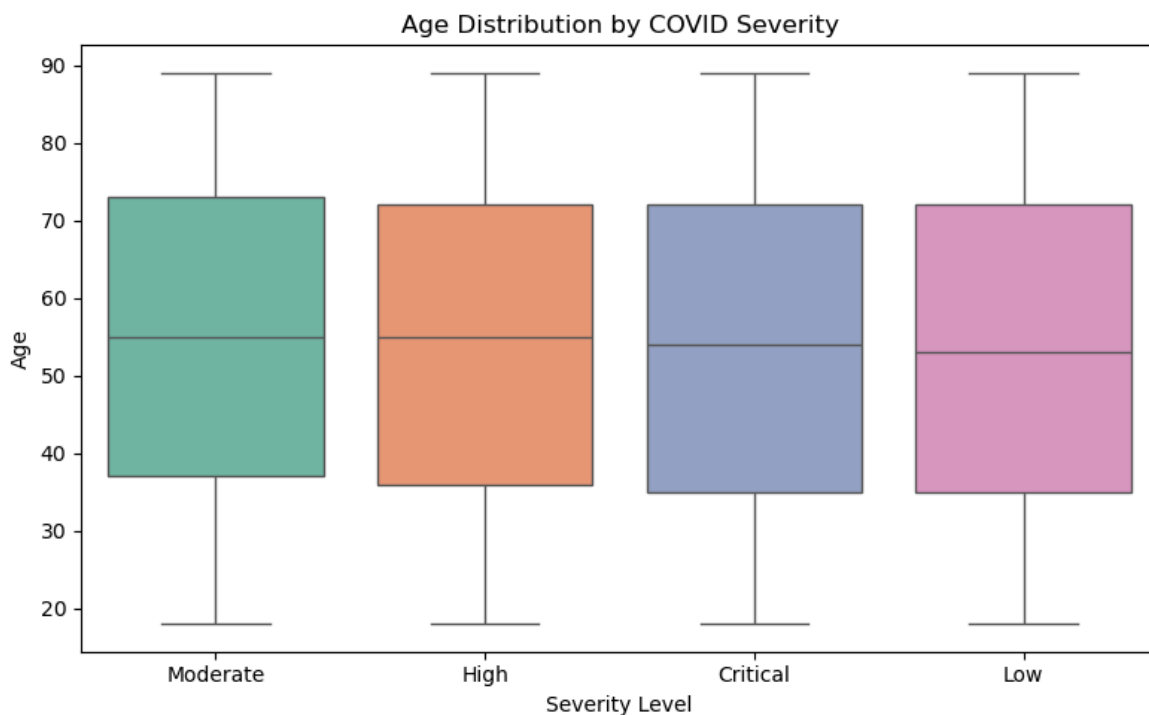
## Age Distribution by Covid Severity

In [213...

```
# Does Age affect COVID Severity?
```

```
plt.figure(figsize=(8, 5))
sns.boxplot(x='Severity', y='Age', data=df, palette='Set2')
plt.title('Age Distribution by COVID Severity')
plt.xlabel('Severity Level')
plt.ylabel('Age')
plt.tight_layout()
plt.show()
```

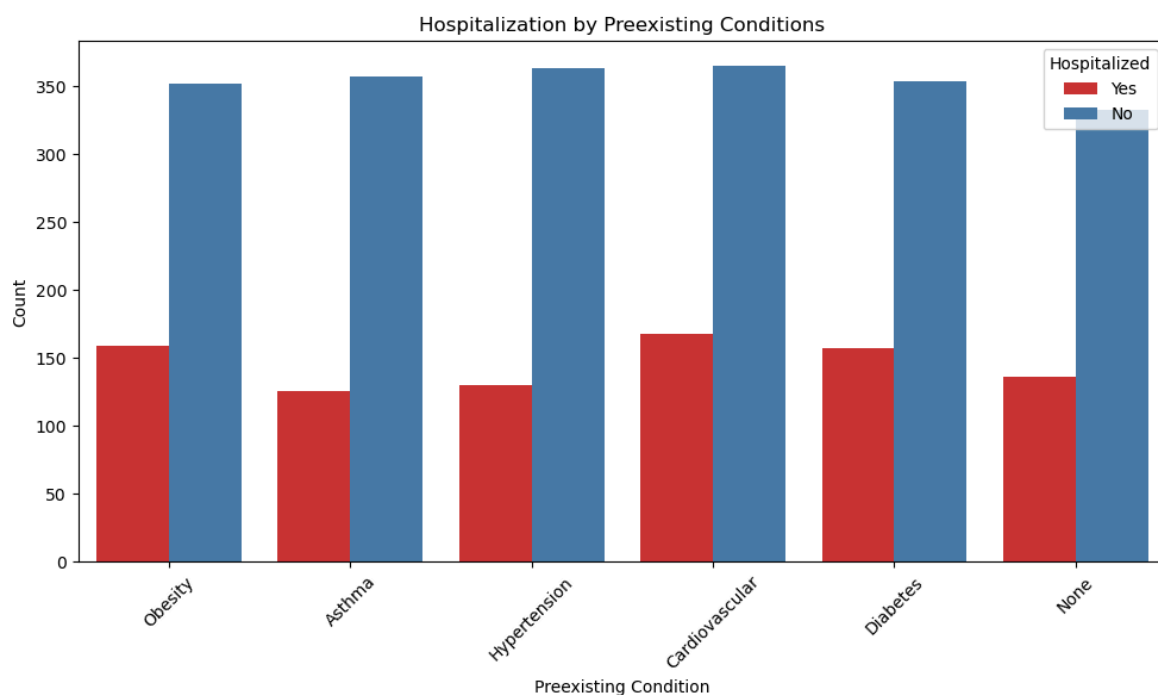




## Hospitalization by Preexisting Conditions


In [214...

```
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='Preexisting_Condition', hue='Hospitalized', palette='S
plt.xticks(rotation=45)
plt.title('Hospitalization by Preexisting Conditions')
plt.xlabel('Preexisting Condition')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```



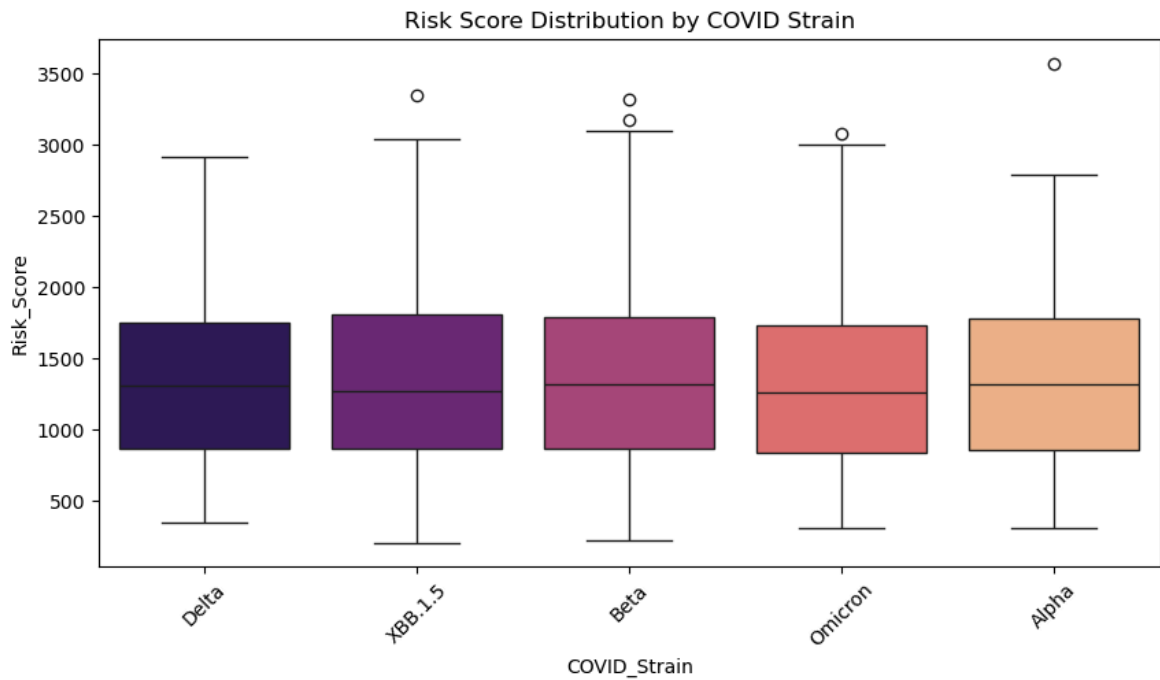
## Gender vs Death Outcome

```
In [215... sns.countplot(data=df, x='Gender', hue='Death', palette=["#ff80ff", "#8000ff"])
plt.title('Death Count by Gender')
plt.ylabel('Number of Patients')
plt.xlabel('Gender')
plt.show()
```



## COVID Strain vs Risk Score

```
In [216... plt.figure(figsize=(10, 5))
sns.boxplot(data=df, x='COVID_Strain', y='Risk_Score', palette='magma')
plt.title('Risk Score Distribution by COVID Strain')
plt.xticks(rotation=45)
plt.show()
```

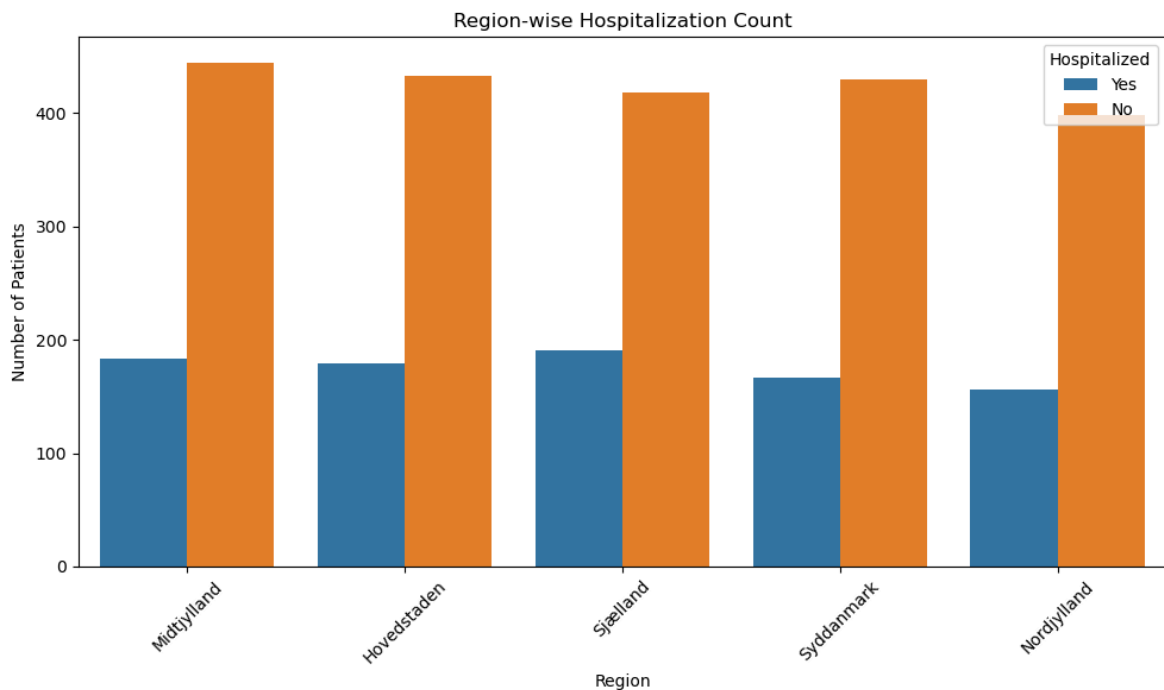


## Region vs Hospitalized

In [217...

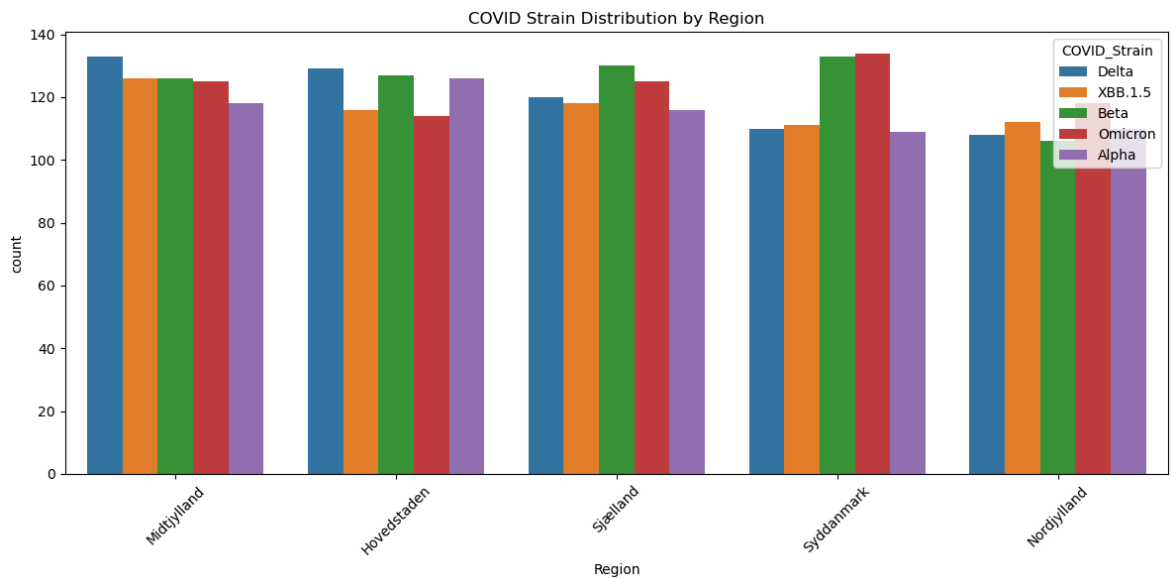
```
plt.figure(figsize=(10,6))
sns.countplot(data=df, x="Region", hue="Hospitalized", order=df["Region"].value_

plt.title("Region-wise Hospitalization Count")
plt.xlabel("Region")
plt.ylabel("Number of Patients")
plt.xticks(rotation=45)
plt.legend(title="Hospitalized")
plt.tight_layout()
plt.show()
```



## Region + COVID Strain vs Hospitalized

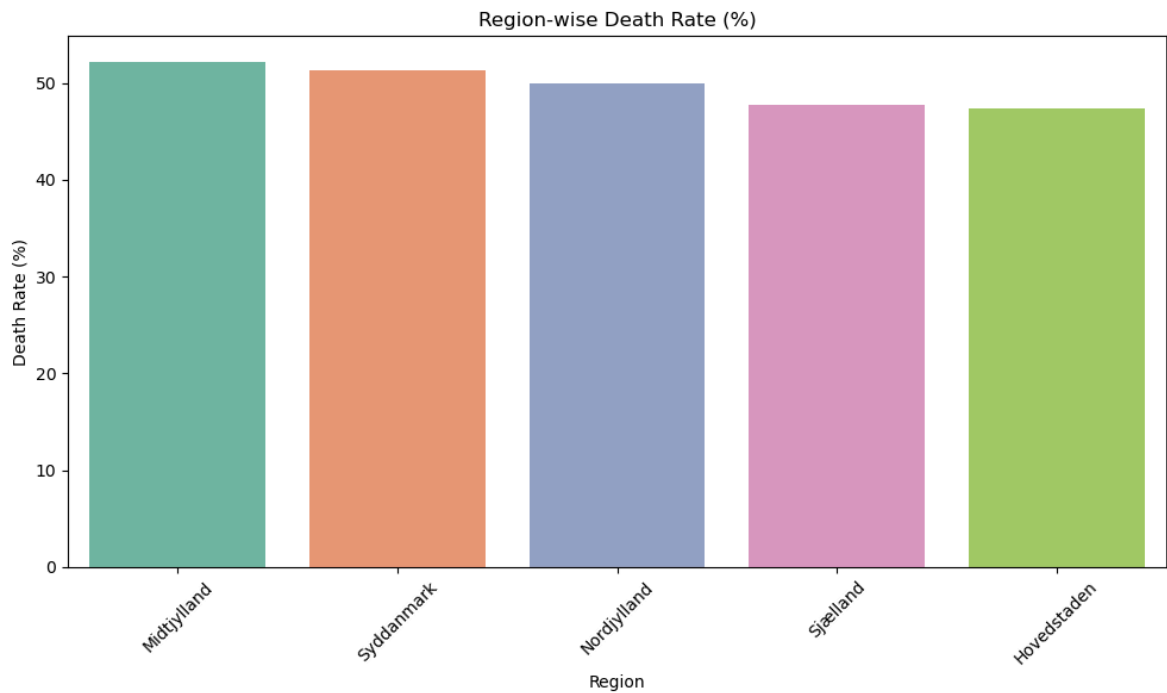
```
In [218... plt.figure(figsize=(12,6))
sns.countplot(data=df, x="Region", hue="COVID_Strain", order=df["Region"].value_
plt.title("COVID Strain Distribution by Region")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



## Region-wise Death Rate(%)

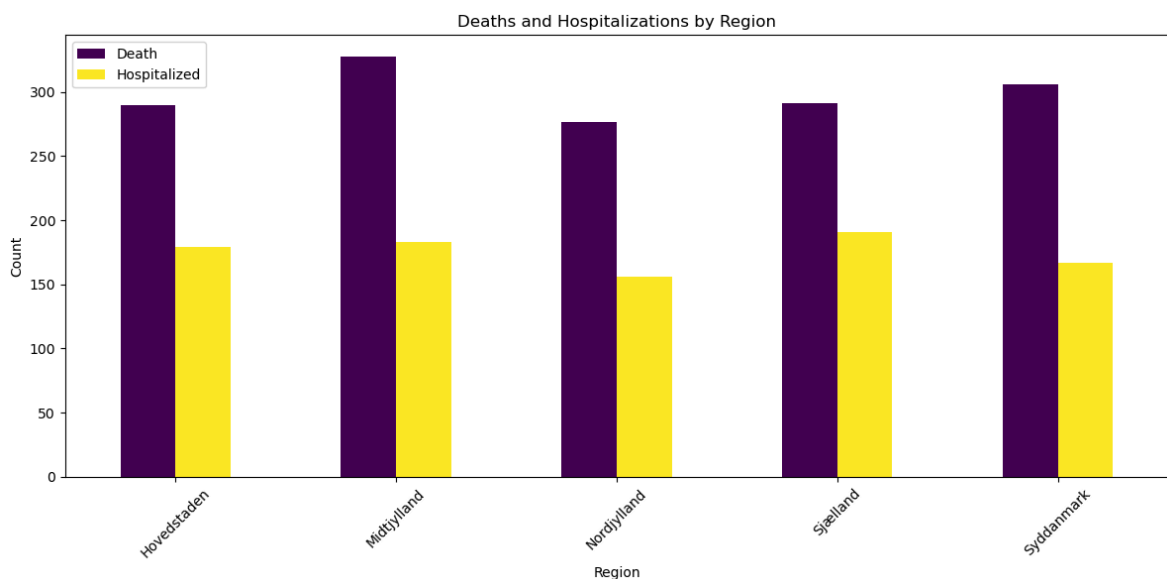
```
In [219... region_death_rate = (
    df.groupby("Region")["Death"]
      .value_counts(normalize=True)
      .unstack()
      .fillna(0) * 100
    ).reset_index()

plt.figure(figsize=(10,6))
sns.barplot(data=region_death_rate, x="Region", y="Yes", order=region_death_rate
plt.title("Region-wise Death Rate (%)")
plt.ylabel("Death Rate (%)")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Which region has the highest number of deaths and hospitalizations?

```
In [220... region_death_hosp = df.groupby('Region')[['Death', 'Hospitalized']].apply(lambda
region_death_hosp.plot(kind='bar', figsize=(12, 6), colormap='viridis')
plt.title('Deaths and Hospitalizations by Region')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```

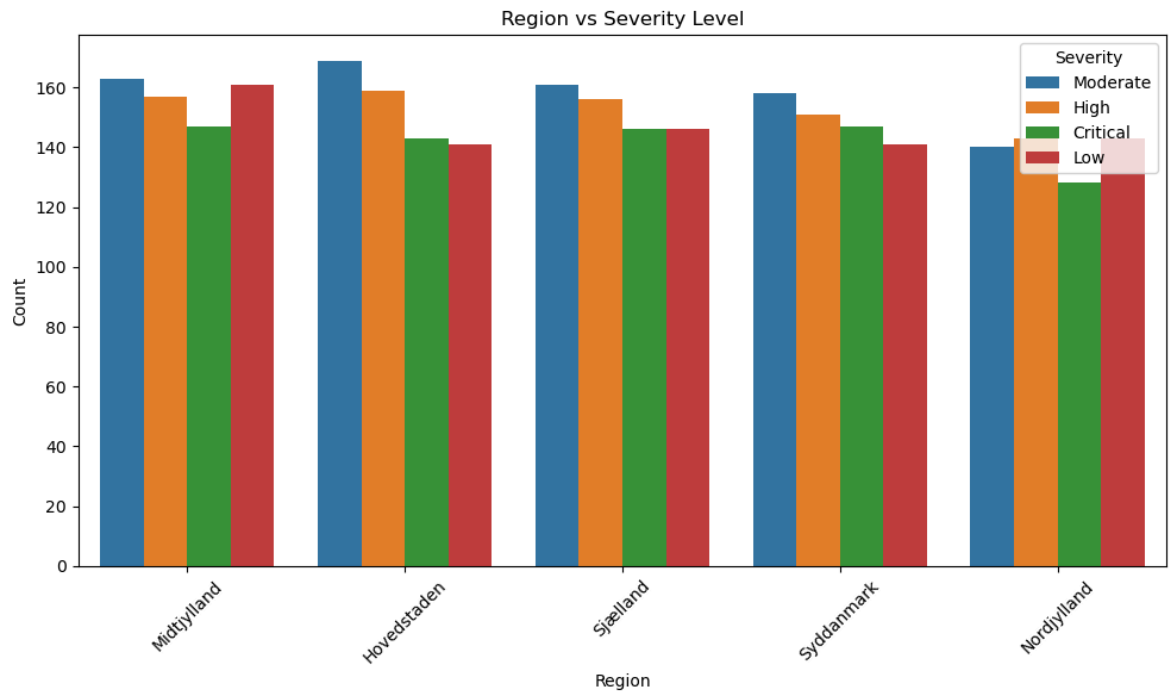


Region vs Severity

```
In [221... plt.figure(figsize=(10,6))
sns.countplot(data=df, x="Region", hue="Severity", order=df["Region"].value_count

plt.title("Region vs Severity Level")
plt.xlabel("Region")
plt.ylabel("Count")
plt.xticks(rotation=45)
```

```
plt.legend(title="Severity")
plt.tight_layout()
plt.show()
```

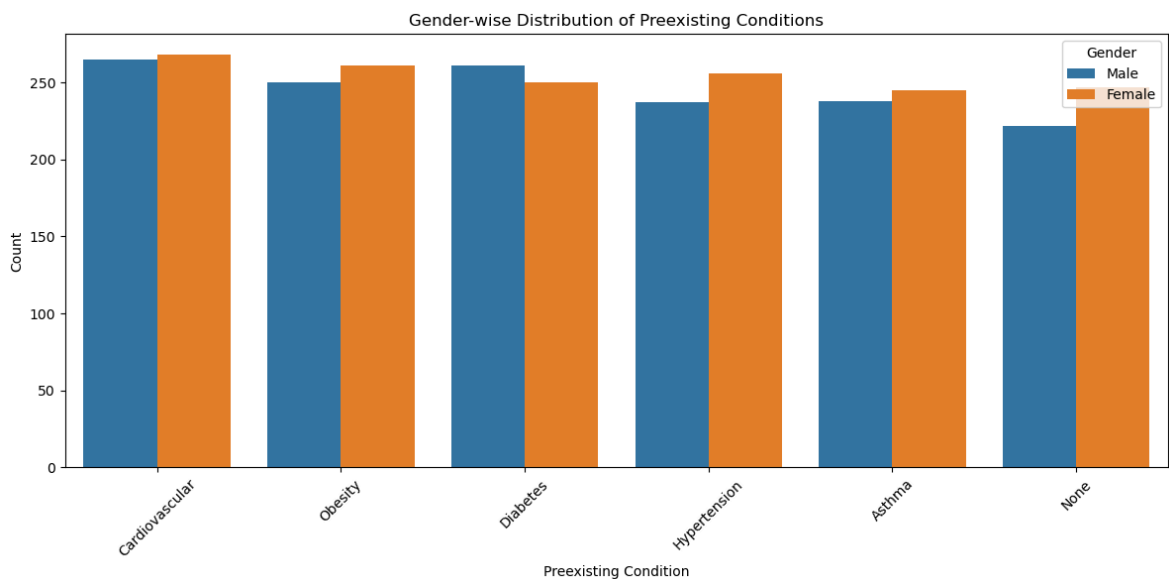


## Gender vs Preexisting\_Condition

In [222...

```
plt.figure(figsize=(12,6))
sns.countplot(data=df, x="Preexisting_Condition", hue="Gender", order=df["Prexi

plt.title("Gender-wise Distribution of Preexisting Conditions")
plt.xlabel("Preexisting Condition")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.legend(title="Gender")
plt.tight_layout()
plt.show()
```

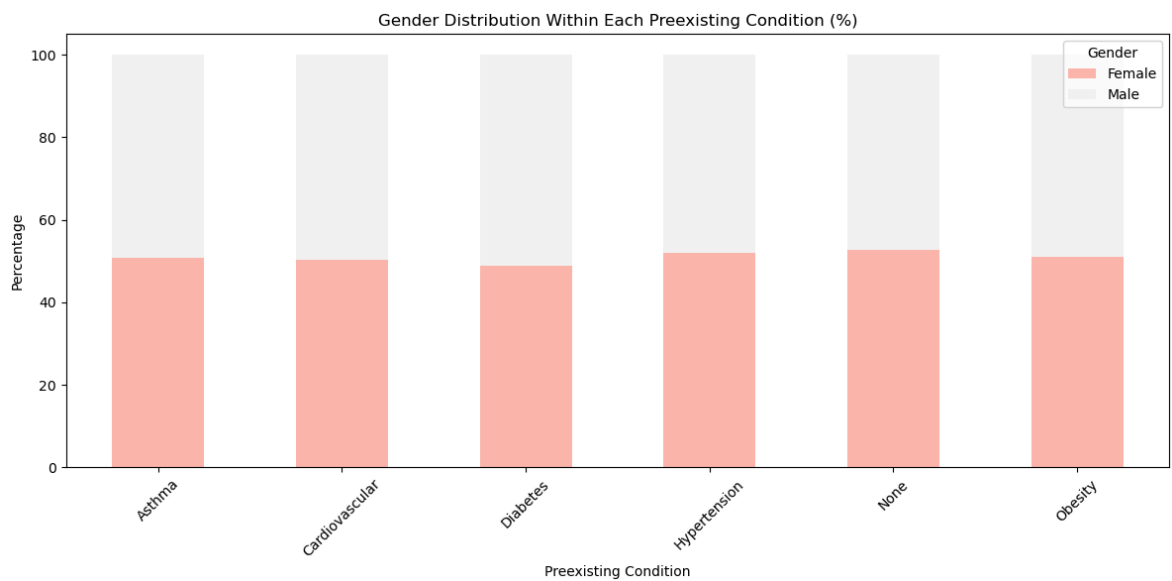


## Gender vs Preexisting\_Condition(%)

In [223...

```
# Create a normalized crosstab
gender_condition_pct = (
    pd.crosstab(df["Preexisting_Condition"], df["Gender"], normalize="index") *
).round(1)

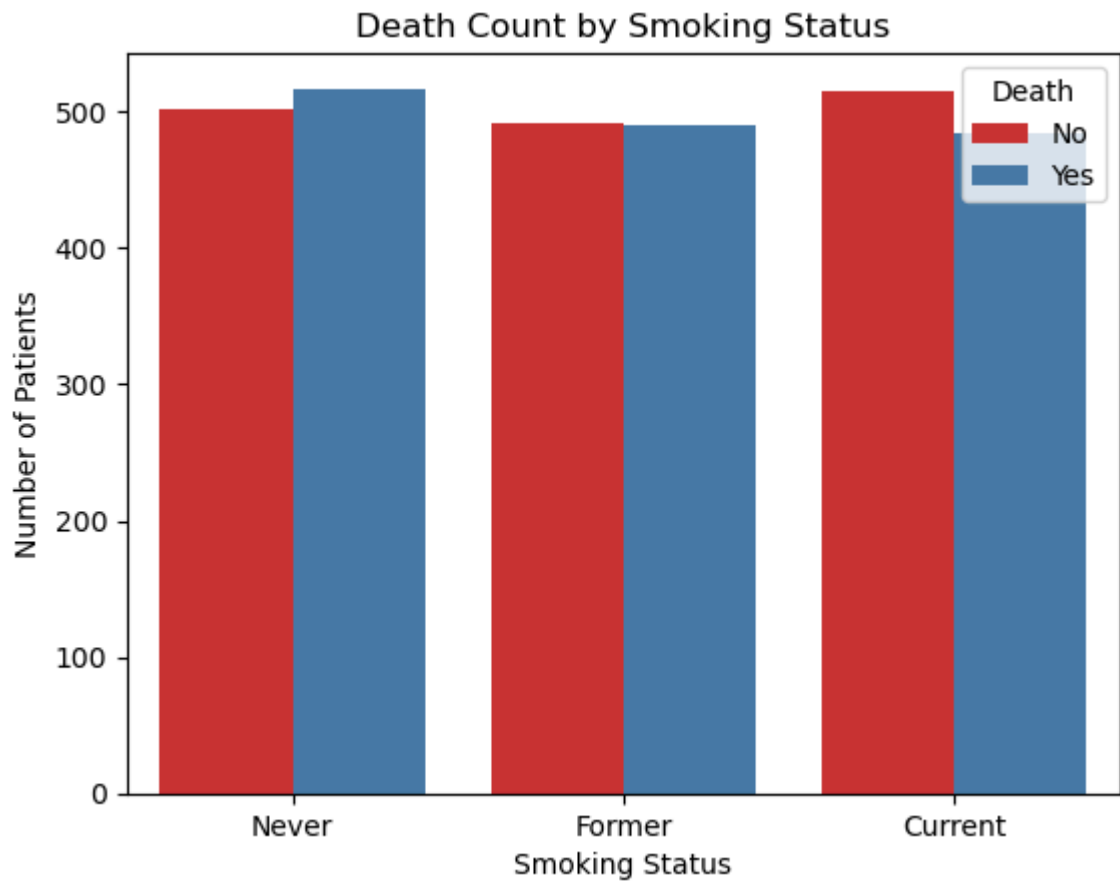
# Plot
gender_condition_pct.plot(kind="bar", stacked=True, figsize=(12,6), colormap="Pa
plt.title("Gender Distribution Within Each Preexisting Condition (%)")
plt.xlabel("Preexisting Condition")
plt.ylabel("Percentage")
plt.xticks(rotation=45)
plt.legend(title="Gender")
plt.tight_layout()
plt.show()
```



## Count Plot: Smoking\_Status vs Death

In [224...

```
sns.countplot(data=df, x='Smoking_Status', hue='Death', palette='Set1')
plt.title('Death Count by Smoking Status')
plt.xlabel('Smoking Status')
plt.ylabel('Number of Patients')
plt.show()
```




### Box Plot: Smoking\_Status vs Death

In [225...

```
sns.boxplot(data=df, x='Smoking_Status', y='BMI', palette='pastel')  
plt.title('BMI Distribution by Smoking Status')  
plt.show()
```





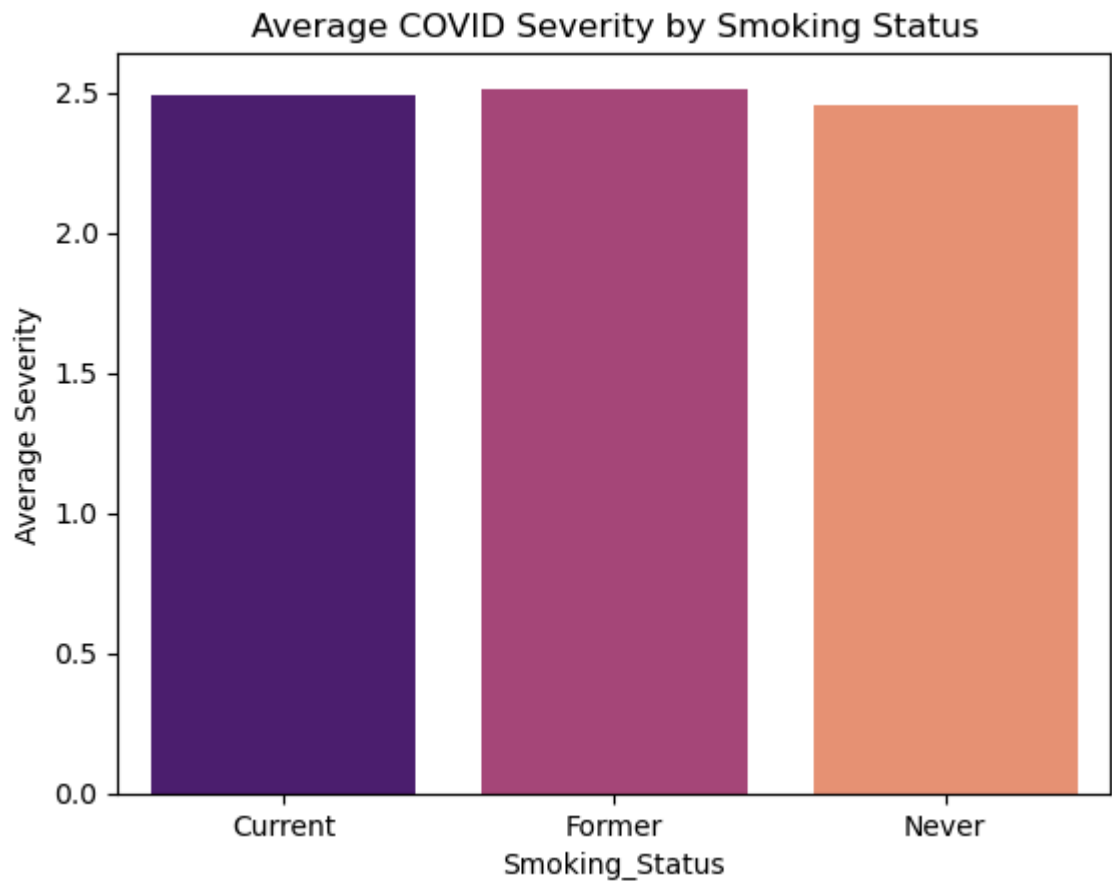
### Box Plot: Average Severity\_Num by Smoking\_Status

In [226...

```
severity_map = {'Low': 1, 'Moderate': 2, 'High': 3, 'Critical': 4}
df['Severity_Num'] = df['Severity'].map(severity_map)

avg_severity = df.groupby('Smoking_Status')['Severity_Num'].mean().reset_index()


sns.barplot(data=avg_severity, x='Smoking_Status', y='Severity_Num', palette='magma')
plt.title('Average COVID Severity by Smoking Status')
plt.ylabel('Average Severity')
plt.show()
```



## Violin Plot: Severity\_Num vs Smoking\_Status


In [227...

```
plt.figure(figsize=(8, 5))
sns.violinplot(data=df, x='Smoking_Status', y='Severity_Num', palette='coolwarm')
plt.title('COVID Severity by Smoking Status')
plt.xlabel('Smoking Status')
plt.ylabel('Severity (Numeric)')
plt.tight_layout()
plt.show()
```



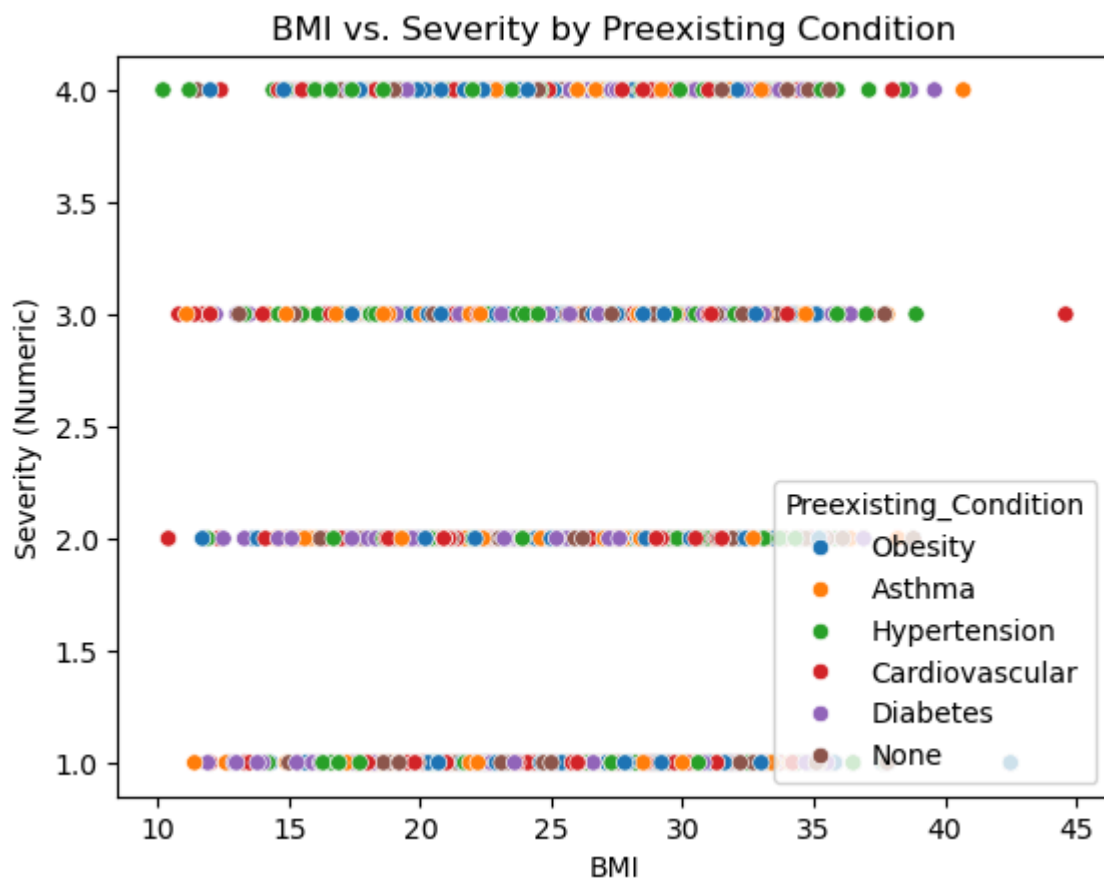
## Age vs. BMI

```
In [228...] sns.scatterplot(data=df, x='Age', y='BMI', hue='Gender')
plt.title('Age vs. BMI by Gender')
plt.show()
```




```
In [229...] sns.scatterplot(data=df, x='BMI', y='Severity_Num', hue='Preexisting_Condition')
plt.title('BMI vs. Severity by Preexisting Condition')
```

```
plt.ylabel('Severity (Numeric)')
plt.show()
```




```
In [230... sns.scatterplot(data=df, x='Age', y='Severity_Num', hue='Hospitalized')
plt.title('Age vs. Severity by Hospitalization')
plt.ylabel('Severity (Numeric)')
plt.show()
```



## Violin Plot: Severity\_Num vs Preexisting\_Condition

In [231...

```
plt.figure(figsize=(12, 6))
sns.violinplot(data=df, x='Preexisting_Condition', y='Severity_Num', palette='Se
plt.title('COVID Severity by Preexisting Condition')
plt.xlabel('Preexisting Condition')
plt.ylabel('Severity (Numeric)')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```




The violin plot shows the distribution of COVID severity scores across different preexisting conditions. Most conditions, including obesity, asthma, hypertension, cardiovascular issues, and diabetes, have a similar median severity level around 2.5–3. However, hypertension and diabetes appear to have a slightly higher range of severity, indicating a tendency toward more severe outcomes. Individuals with no preexisting conditions generally show a narrower severity spread, suggesting milder cases on average. Overall, the plot suggests that while preexisting conditions may impact severity, the variation within each group is still quite broad.

## Plots on the basis of Occupation

### Bar Plot: Average Severity by Occupation


```
In [232... avg_severity = df.groupby('Occupation')['Severity_Num'].mean().reset_index()

plt.figure(figsize=(12,6))
sns.barplot(data=avg_severity, x='Occupation', y='Severity_Num', palette='viridi
plt.title('Average COVID Severity by Occupation')
plt.xticks(rotation=45)
plt.ylabel('Avg Severity (1=Low, 4=Critical)')
plt.tight_layout()
plt.show()
```



### Count Plot: Hospitalization by Occupation


```
In [233... plt.figure(figsize=(12,6))
sns.countplot(data=df, x='Occupation', hue='Hospitalized', palette=["#008080", "#FFD700"])
plt.title('Hospitalization Count by Occupation')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



## Count Plot: Number of Patients by Occupation

In [234...

```
plt.figure(figsize=(12,6))
sns.countplot(data=df, x='Occupation', order=df['Occupation'].value_counts().index)
plt.title('Number of Patients by Occupation')
plt.xticks(rotation=45)
plt.show()
```



Most patients are from Healthcare, Office Worker, and Student groups.


Very few patients belong to Retired or Teacher categories, indicating different exposure or age distribution.

## Stacked Bar Plot: Death Count by Occupation

In [235...

```
death_by_occ = df.groupby(['Occupation', 'Death']).size().unstack().fillna(0)

death_by_occ.plot(kind='bar', stacked=True, figsize=(12,6), colormap='coolwarm')
plt.title('Death by Occupation')
plt.ylabel('Number of Patients')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



## Multivariate Plots

### Pie Charts: Gender vs Symptoms

In [236...

```
# Define new unique colors
symptom_colors = {
    "Mild": "#9b59b6",
    "Moderate": "#800020",
    "Severe": "#004d4d"
}

# Male Pie Chart
male_symptoms = df[df["Gender"] == "Male"]["Symptoms"].value_counts()
male_colors = [symptom_colors[s] for s in male_symptoms.index]

plt.figure(figsize=(5,5))
plt.pie(male_symptoms, labels=male_symptoms.index, autopct='%1.1f%%',
        colors=male_colors, startangle=90)
plt.title("Male Symptom Distribution")
plt.tight_layout()
plt.show()


# Female Pie Chart
female_symptoms = df[df["Gender"] == "Female"]["Symptoms"].value_counts()
female_colors = [symptom_colors[s] for s in female_symptoms.index]

plt.figure(figsize=(5,5))
plt.pie(female_symptoms, labels=female_symptoms.index, autopct='%1.1f%%',
        colors=female_colors, startangle=90)
plt.title("Female Symptom Distribution")
```




```
plt.tight_layout()  
plt.show()
```

Male Symptom Distribution




## Female Symptom Distribution




## BMI Category vs Severity

```
In [237... crosstab = pd.crosstab(df['BMI_Category'], df['Severity'])
sns.heatmap(crosstab, annot=True, cmap='YlGnBu', fmt='d')
plt.title('Severity by BMI Category')
plt.xlabel('Severity')
plt.ylabel('BMI Category')
plt.show()
```




Is there any relationship between vaccine doses and recovery time?

```
In [238... plt.figure(figsize=(10, 6))
sns.scatterplot(data=df, x='Doses_Received', y='Recovery_Days', hue='Severity',
plt.title('Doses Received vs. Days of Recovery')
plt.tight_layout()
plt.show()
```



```
In [239... sns.catplot(x='Smoking_Status', y='Recovery_Days', kind='swarm', data=df)
plt.title("Average Recovery Days by Smoking Status")
```


```
plt.xticks(rotation=45)
plt.show()
```



## Analysis Questions

1. Which region has the highest number of deaths and hospitalizations?

```
In [241... region_death_hosp = df.groupby('Region')[['Death', 'Hospitalized']].apply(lambda
region_death_hosp.plot(kind='bar', figsize=(12, 6), colormap='viridis')
plt.title('Deaths and Hospitalizations by Region')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



## 2. Create Binary Tables

```
In [242... # Create binary column: 1 if Doses_Received is a number, 0 if 'None' or invalid
df['Doses_Received_Binary'] = pd.to_numeric(df['Doses_Received'], errors='coerce')

# Create binary column: 1 if Recovery_Days is a number, 0 if 'None' or invalid
df['Recovery_Days_Binary'] = pd.to_numeric(df['Recovery_Days'], errors='coerce')
```

```
In [243... print(df[['Doses_Received', 'Doses_Received_Binary']].head())
```

	Doses_Received	Doses_Received_Binary
0	1	1
1	0	1
2	3	1
3	1	1
4	2	1

```
In [247... print(df[['Recovery_Days', 'Recovery_Days_Binary']].head())
```

	Recovery_Days	Recovery_Days_Binary
0	302	1
1	0	0
2	0	0
3	546	1
4	0	0

## 3. Is there a trend between vaccine doses and recovery days?

```
In [248... correlation = df[['Doses_Received_Binary', 'Recovery_Days_Binary']].corr().loc['Doses_Received_Binary', 'Recovery_Days_Binary']
print(f"Correlation between Doses Received and Days to Recovery: {correlation}")
```

Correlation between Doses Received and Days to Recovery: nan

## 4. Which region has the highest hospitalization and death count?

```
In [249... # Count of 'Yes' in Hospitalized and Death columns per region
region_stats = df.groupby('Region')[['Hospitalized', 'Death']].apply(lambda x: (
print(region_stats.sort_values(by='Death', ascending=False))
```

	Hospitalized	Death
Region		
Midtjylland	183	328
Syddanmark	167	306
Sjælland	191	291
Hovedstaden	179	290
Nordjylland	156	277

## 5. What occupations are most associated with high severity and death?

```
In [250... # Mean severity and death count per occupation
severity_by_occ = df.groupby('Occupation')['Severity_Num'].mean().sort_values(ascending=True)
death_by_occ = df[df['Death'] == 'Yes'].groupby('Occupation').size().sort_values(ascending=True)
print("Severity by Occupation:\n", severity_by_occ)
print("\nDeath Count by Occupation:\n", death_by_occ)
```

```
Severity by Occupation:
Occupation
Office Worker    2.548708
Unemployed       2.507042
Teacher          2.496945
Driver           2.483740
Student          2.472795
Healthcare       2.400826
Name: Severity_Num, dtype: float64
```

```
Death Count by Occupation:
Occupation
Unemployed       265
Office Worker    253
Student          251
Teacher          245
Healthcare       241
Driver           237
dtype: int64
```

## 6. Are elderly people more at risk of death and hospitalization?

```
In [251... # Create age group bins
df['Age_Group'] = pd.cut(df['Age'], bins=[0, 18, 40, 60, 80, 100],
                        labels=['0-18', '19-40', '41-60', '61-80', '81+'])

risk_by_age = df.groupby('Age_Group')[['Hospitalized', 'Death']].apply(lambda x:
print(risk_by_age)
```

	Hospitalized	Death
Age_Group		
0-18	17	19
19-40	255	460
41-60	246	408
61-80	250	429
81+	108	176

## 7. Which type of patients took the longest time to recover?

```
In [252... long_recovery = df.groupby('Preexisting_Condition')['Recovery_Days'].mean().round(2)
print(long_recovery)
```

```

Preexisting_Condition
Diabetes      195
Asthma        190
Cardiovascular 185
Obesity       182
Hypertension  173
None          168
Name: Recovery_Days, dtype: int32

```

## 8. Which patients with preexisting conditions experienced the highest severity?

```

In [253... severity_condition = df.groupby('Preexisting_Condition')['Severity_Num'].mean().
print(severity_condition)

```

```

Preexisting_Condition
Hypertension      3
Asthma            2
Cardiovascular    2
Diabetes          2
None              2
Obesity           2
Name: Severity_Num, dtype: int32

```

## 9. Which region had the highest number of hospitalizations (possibly indicating lower healthcare access)?

```

In [254... hosp_by_region = df.groupby('Region')['Hospitalized'].apply(lambda x: (x == 'Yes'
print(hosp_by_region)

```

```

Region
Sjælland      191
Midtjylland   183
Hovedstaden   179
Syddanmark    167
Nordjylland   156
Name: Hospitalized, dtype: int64

```

## 10. Which age group had the longest recovery time?

```

In [255... df['Age_Group'] = pd.cut(df['Age'], bins=[0,18,40,60,80,100], labels=['0-18','19
recovery_by_age = df.groupby('Age_Group')['Recovery_Days'].mean().round().astype
print(recovery_by_age)

```

```

Age_Group
0-18      240
81+       205
19-40     187
61-80     176
41-60     169
Name: Recovery_Days, dtype: int32

```

## 11. Did patients with fewer vaccine doses have a longer recovery time?

```

In [256... # Group by doses and calculate mean recovery time
recovery_by_dose = df.groupby('Doses_Received')['Recovery_Days'].mean().round().
print("Average Recovery Time by Number of Doses:\n")
print(recovery_by_dose)

```

Average Recovery Time by Number of Doses:

```
Doses_Received
0    190
1    172
2    161
3    188
Name: Recovery_Days, dtype: int32
```

```
In [257... correlation = df[['Doses_Received_Binary', 'Recovery_Days_Binary']].corr().loc['
print("Correlation between Doses Received and Recovery Days :", correlation)
```

Correlation between Doses Received and Recovery Days : nan

## 12. Which gender had a higher number of deaths or severe cases?

```
In [258... gender_severity = df.groupby('Gender')['Severity_Num'].mean()
death_gender = df.groupby('Gender')['Death'].apply(lambda x: (x == 'Yes').sum())
print("Severity:\n", gender_severity)
print("\nDeath Count:\n", death_gender)
```

```
Severity:
Gender
Female    2.491814
Male      2.478615
Name: Severity_Num, dtype: float64
```

```
Death Count:
Gender
Female    784
Male      708
Name: Death, dtype: int64
```

## 13. Which occupation group had a higher severity or death rate?

```
In [259... occ_severity = df.groupby('Occupation')['Severity_Num'].mean().sort_values(ascen
occ_death = df[df['Death'] == 'Yes'].groupby('Occupation').size().sort_values(as
print("Severity by Occupation:\n", occ_severity)
print("\nDeath Count by Occupation:\n", occ_death)
```

```
Severity by Occupation:
Occupation
Office Worker    2.548708
Unemployed       2.507042
Teacher          2.496945
Driver           2.483740
Student          2.472795
Healthcare       2.400826
Name: Severity_Num, dtype: float64
```

```
Death Count by Occupation:
Occupation
Unemployed       265
Office Worker    253
Student          251
Teacher          245
Healthcare       241
Driver           237
dtype: int64
```



## 14. Does COVID severity tend to be higher among smokers?

```
In [260... severity_by_smoke = df.groupby('Smoking_Status')['Severity_Num'].mean().sort_val
print(severity_by_smoke)
```

```
Smoking_Status
Former      2.513761
Current     2.492000
Never       2.451423
Name: Severity_Num, dtype: float64
```

## 15. Which smoking group had the longest recovery time?

```
In [261... # check average recovery time per smoking group
recovery_by_smoke = df.groupby('Smoking_Status')['Recovery_Days'].mean().round(0)
print("Average Recovery Days by Smoking Status (rounded):\n", recovery_by_smoke)
```

```
Average Recovery Days by Smoking Status (rounded):
Smoking_Status
Current      186
Never        184
Former       176
Name: Recovery_Days, dtype: int32
```

## 16. Which BMI group had the longest recovery time?

```
In [262... # Use proper recovery days column
recovery_by_bmi = df.groupby('BMI_Category')['Recovery_Days'].mean().round(0).as
print("Average Recovery Days by BMI Category:\n", recovery_by_bmi)
```

```
Average Recovery Days by BMI Category:
BMI_Category
Underweight  193
Normal       183
Overweight   181
Obese        177
Name: Recovery_Days, dtype: int32
```

## 17. Does the combination of smoking and obesity increase the risk of death?

```
In [263... combo_death = df.groupby(['Smoking_Status', 'BMI_Category'])['Death'].apply(lamb
print(combo_death.sort_values(ascending=False))
```

```
Smoking_Status BMI_Category
Never          Normal      202
Current        Normal      196
Never          Overweight   186
Former         Normal      177
               Overweight   175
Current        Overweight   164
Former         Obese        92
Current        Obese        84
Never          Obese        76
               Underweight   53
Former         Underweight   46
Current        Underweight   41
Name: Death, dtype: int64
```

## 18. Which combination of smoking and preexisting condition is associated with the highest severity?

```
In [264... combo_severity = df.groupby(['Smoking_Status', 'Preexisting_Condition'])['Severity'].head(10)
```

Smoking_Status	Preexisting_Condition	Severity
Current	Asthma	3
Former	Cardiovascular	3
Never	Hypertension	3
Former	Obesity	3
	Diabetes	3
	Hypertension	3
Current	Hypertension	3
	Diabetes	3
Former	Asthma	2
Current	Obesity	2

Name: Severity\_Num, dtype: int32

## 19. Is there any relationship between BMI and vaccine doses (e.g., do individuals with lower BMI tend to receive fewer doses)?

```
In [265... bmi_dose_corr = df[['BMI', 'Doses_Received']].corr().loc['BMI', 'Doses_Received']
print("Correlation between BMI and Vaccine Doses:", bmi_dose_corr)
```

Correlation between BMI and Vaccine Doses: -0.015218010745538513

## 20. Is the death rate higher among smokers despite receiving vaccine doses?

```
In [266... smoke_death = df.groupby('Smoking_Status')['Death'].apply(lambda x: (x == 'Yes').sum())
print(smoke_death)
```

Smoking_Status	Death
Current	485
Former	490
Never	517

Name: Death, dtype: int64

## 21. Even though obese individuals received more vaccine doses, how was their recovery time?

```
In [270... obese_data = df[df['BMI_Category'].isin(['Obese', 'Severely Obese'])]
obese_recovery = obese_data.groupby('Doses_Received')['Recovery_Days'].mean().round(2)
print(obese_recovery)
```

Doses_Received	Recovery_Days
0	181.0
1	161.0
2	172.0
3	182.0

Name: Recovery\_Days, dtype: float64

## Vaccine Doses Based Data Analysis

### 1. Does taking more vaccine doses reduce the severity of COVID?

Analysis: Check average severity score (Severity\_Num) grouped by Doses\_Received

```
In [275... df.groupby('Doses_Received')['Severity_Num'].mean().round()
```

```
Out[275... Doses_Received
0      2.0
1      3.0
2      2.0
3      3.0
Name: Severity_Num, dtype: float64
```

## 2. Do more vaccine doses reduce the recovery time?

Analysis: Compare Recovery\_Days across vaccine dose groups.

```
In [277... df.groupby('Doses_Received')['Recovery_Days'].mean().round()
```

```
Out[277... Doses_Received
0      190.0
1      172.0
2      161.0
3      188.0
Name: Recovery_Days, dtype: float64
```

## 3. What is the death rate for each vaccine dose group?

Analysis: Check percentage of 'Yes' in Death column for each Doses\_Received.

```
In [284... death_by_dose = df.groupby('Doses_Received')['Death'].value_counts(normalize=True)
print(death_by_dose)
```

```
Doses_Received
0      48.625654
1      50.403226
2      52.192067
3      50.100604
Name: Yes, dtype: float64
```

## 4. How many patients who received 3 doses were hospitalized?

Analysis: Count Hospitalized status among people with 3 doses.

```
In [285... df[(df['Doses_Received'] == 3) & (df['Hospitalized'] == 'Yes')].shape[0]
```

```
Out[285... 143
```

## 5. Which age group received the most vaccine doses (>=2)?

Analysis: Group by Age\_Group and count how many received  $\geq 2$  doses.

```
In [286... df[df['Doses_Received'] >= 2].groupby('Age_Group').size().sort_values(ascending=
```

```
Out[286...] Age_Group
61-80      292
19-40      285
41-60      258
81+        124
0-18        17
dtype: int64
```

## 6. Is there a relationship between vaccine dose count and preexisting conditions?

Analysis: Cross-tabulate Preexisting\_Condition with Doses\_Received.

```
In [287...] pd.crosstab(df['Preexisting_Condition'], df['Doses_Received'])
```

```
Out[287...]
      Doses_Received  0  1  2  3
Preexisting_Condition
Asthma             244  70  90  79
Cardiovascular     276  80  80  97
Diabetes           259  82  70  100
Hypertension       252  86  83  72
None               236  84  80  69
Obesity            261  94  76  80
```

## 7. Among vaccinated people ( $\geq 1$ dose), which preexisting condition had the most deaths?

Analysis: Focus only on vaccinated people and see deaths by condition.

```
In [288...] df_vac = df[df['Doses_Received'] >= 1]
df_vac[df_vac['Death'] == 'Yes']['Preexisting_Condition'].value_counts()
```

```
Out[288...] Preexisting_Condition
Obesity             130
Diabetes            127
Cardiovascular      126
None                126
Hypertension        125
Asthma              115
Name: count, dtype: int64
```

## 8. Is there any effect of vaccine doses on symptoms?

Analysis: Average count of symptoms per dose group.

```
In [289...] df['Symptom_Count'] = df['Symptoms'].fillna('').apply(lambda x: len(x.split(',')))
df.groupby('Doses_Received')['Symptom_Count'].mean()
```

```
Out[289... Doses_Received
0      1.0
1      1.0
2      1.0
3      1.0
Name: Symptom_Count, dtype: float64
```

## 9. Among patients with 0 vaccine doses, what's the most common preexisting condition?

Analysis: Find most common health issue in unvaccinated group.

```
In [290... df[df['Doses_Received'] == 0]['Preexisting_Condition'].value_counts()
```

```
Out[290... Preexisting_Condition
Cardiovascular      276
Obesity              261
Diabetes             259
Hypertension        252
Asthma              244
None                236
Name: count, dtype: int64
```

## 10. Are vaccinated people less likely to be hospitalized?

Analysis: Compare hospitalization rate among 0, 1, 2, 3 dose groups.

```
In [292... df.groupby('Doses_Received')['Hospitalized'].value_counts(normalize=True).unstack()
```

```
Out[292... Doses_Received
0      28.795812
1      29.435484
2      30.688935
3      28.772636
Name: Yes, dtype: float64
```

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]:

In [ ]: